

csci 1470

Eric Ewing

Wednesday,  
4/2/25

# Deep Learning

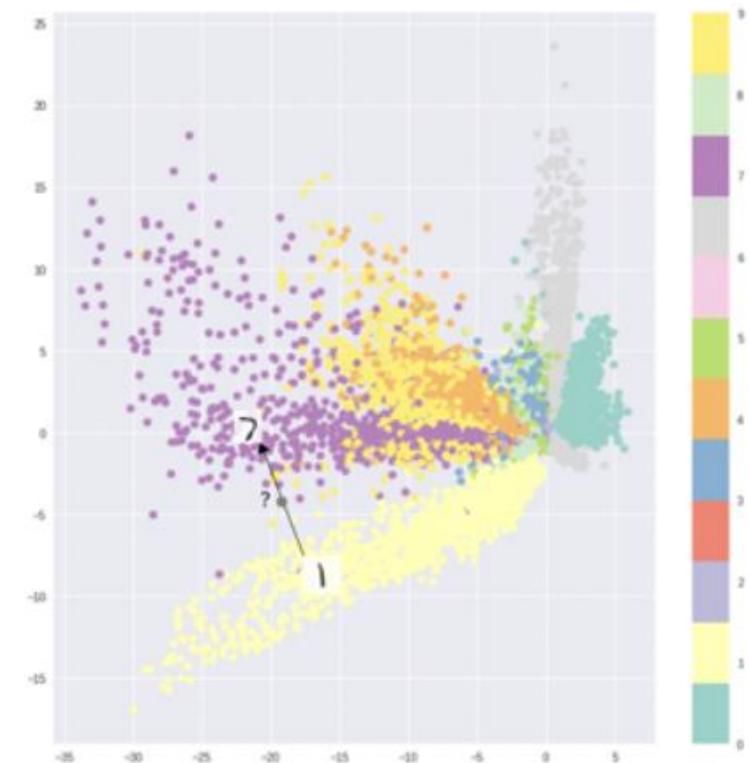
Day 25: Image Generation Day 2: VAEs

# Logistics

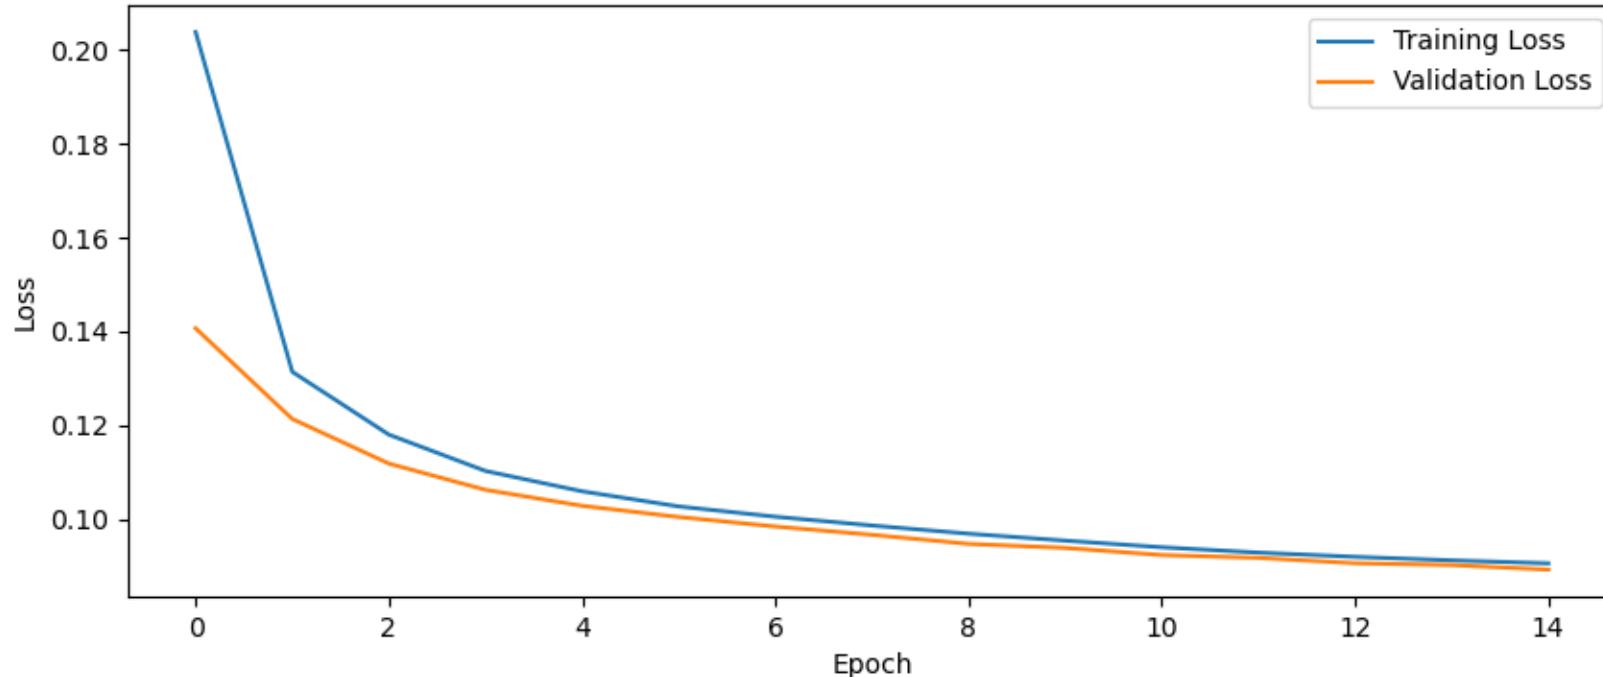
- Weekly Quiz is out
- Make slow and steady progress on your final project!
  - Deep Learning is not something that can all happen at the last minute, data takes time to process, experiments take time to run

# Generating Images

- How can we generate a “new” image using a decoder?
- Sample a vector in latent space and send it to the decoder...
- But how do you choose which vector?
- What if you wanted to generate a specific image? How would you find the right vector?



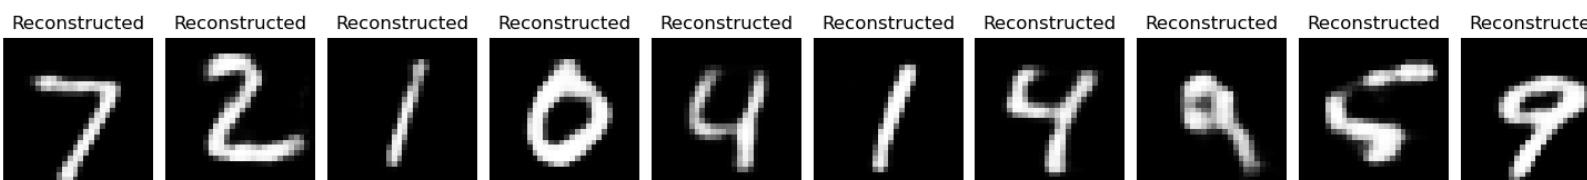
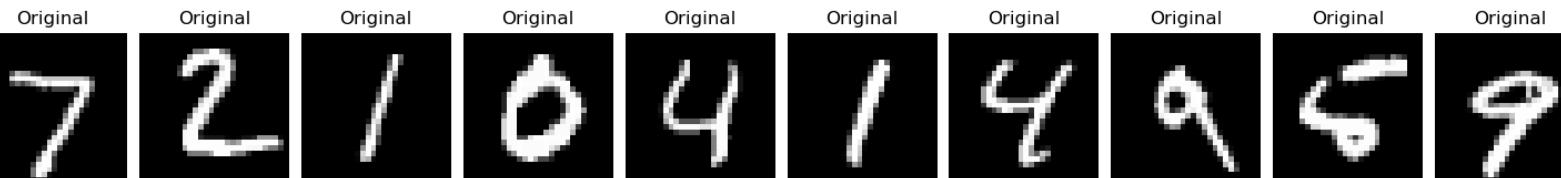
Autoencoder Loss

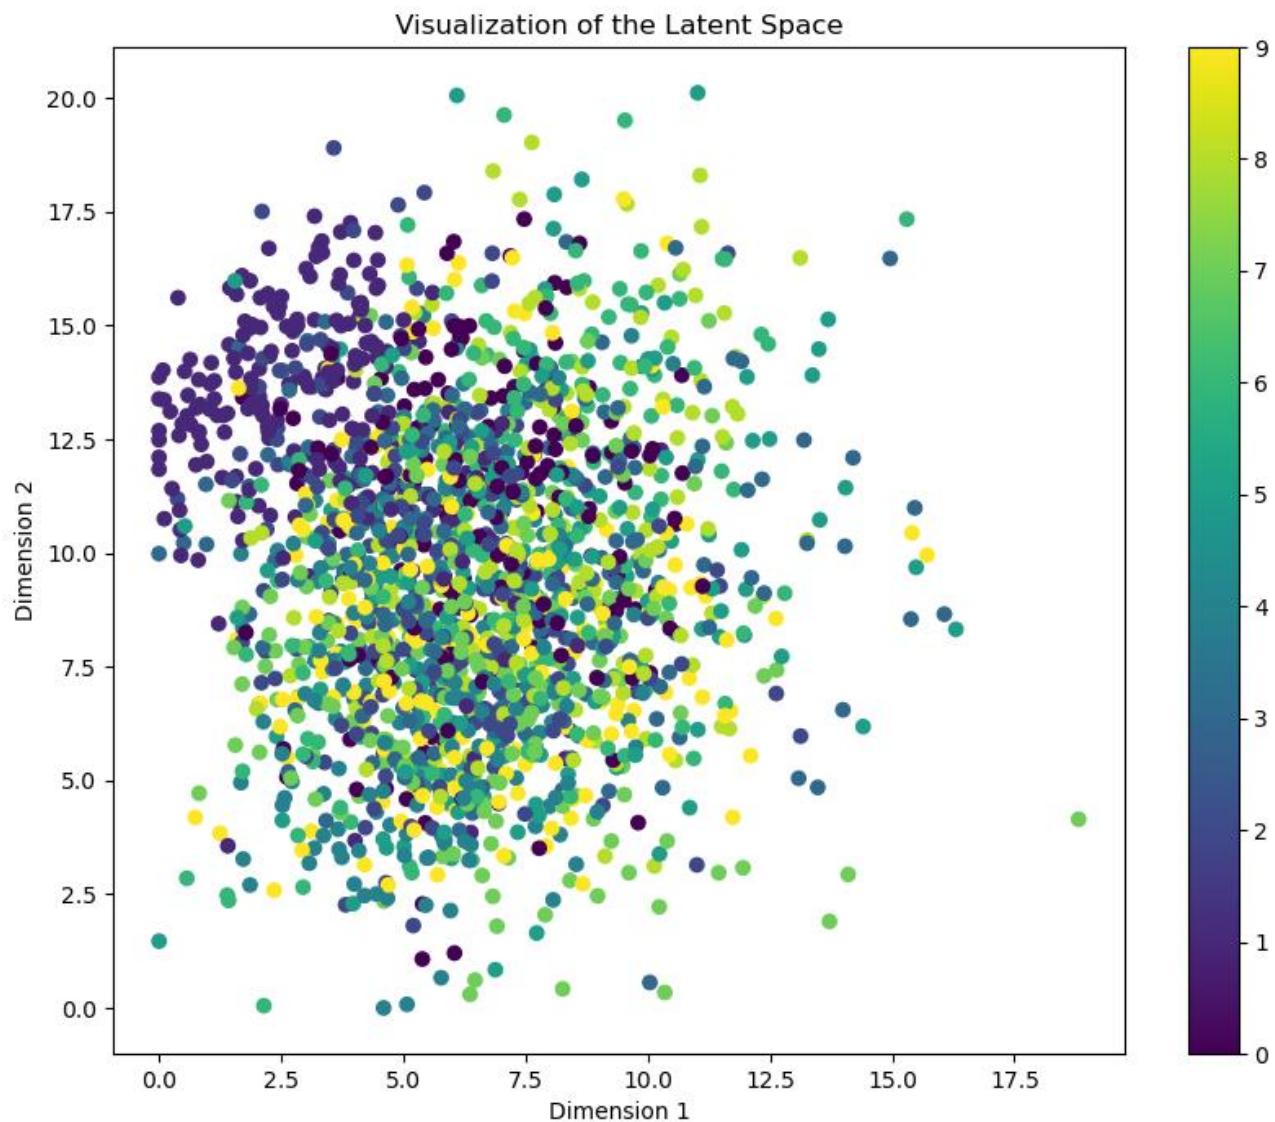


(Encoding size = 32)

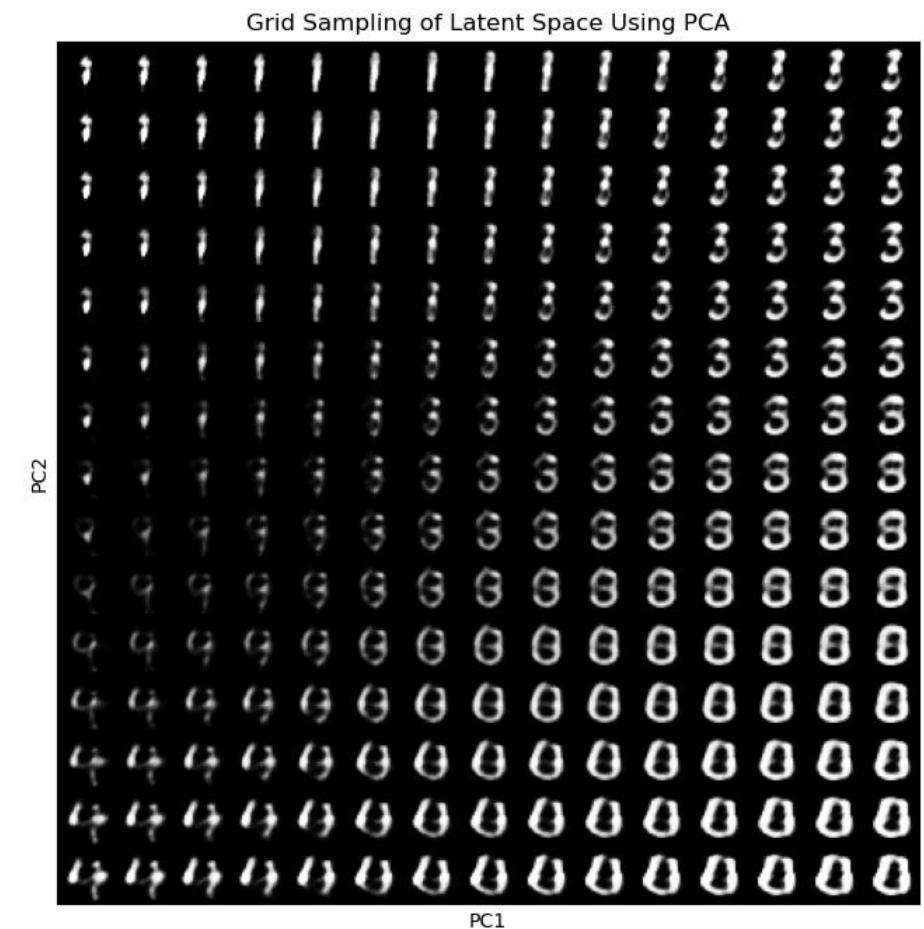
Is this a good  
autoencoder?

Why is loss alone (even  
with validation loss) not  
enough to tell us?





Grid sample latent space and pass to encoder



Explained variance: PC1=0.32, PC2=0.11

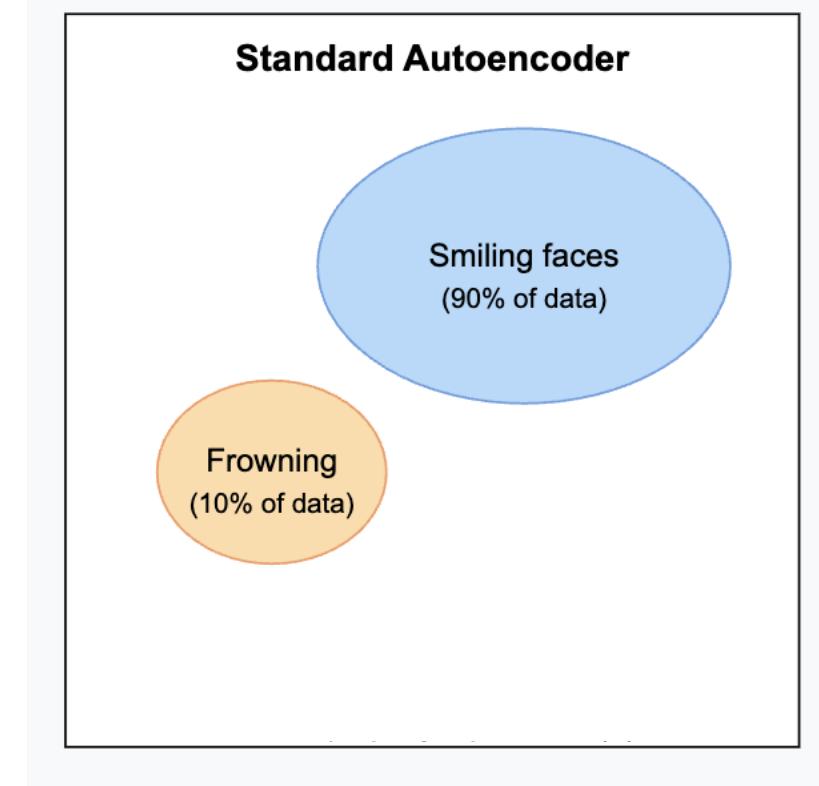
# Autoencoders can generate, but are not generative

Recall:

- Discriminative models learn  $P(y | X)$
- Generative models learn  $P(X)$

When we randomly sample, we may get some “invalid” outputs. A generative model could assign these invalid outputs a low probability  $P(X)$

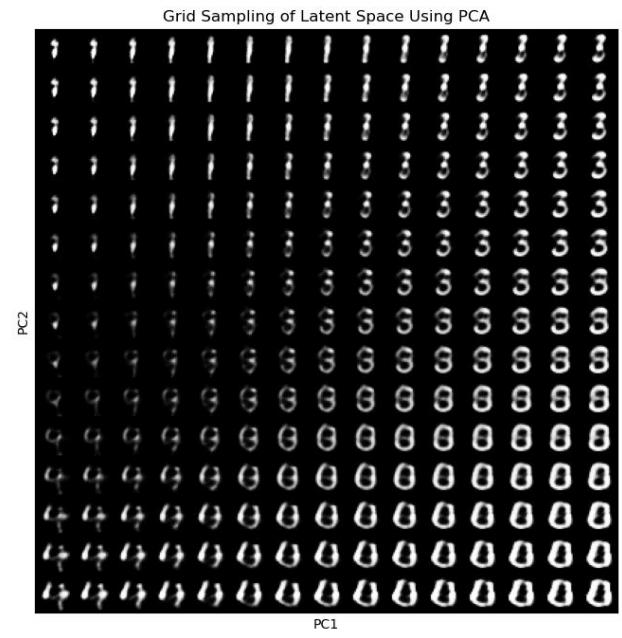
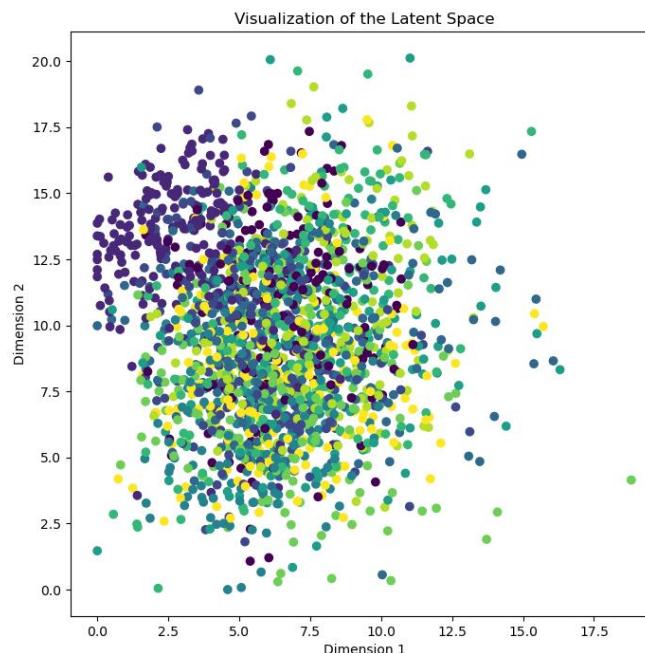
Nothing constrains the latent space of an autoencoder to represent probability distributions



# Issues with Autoencoders

- Vectors close together in latent space may not produce similar outputs
- Tend to overfit data (struggle to produce “new” outputs)

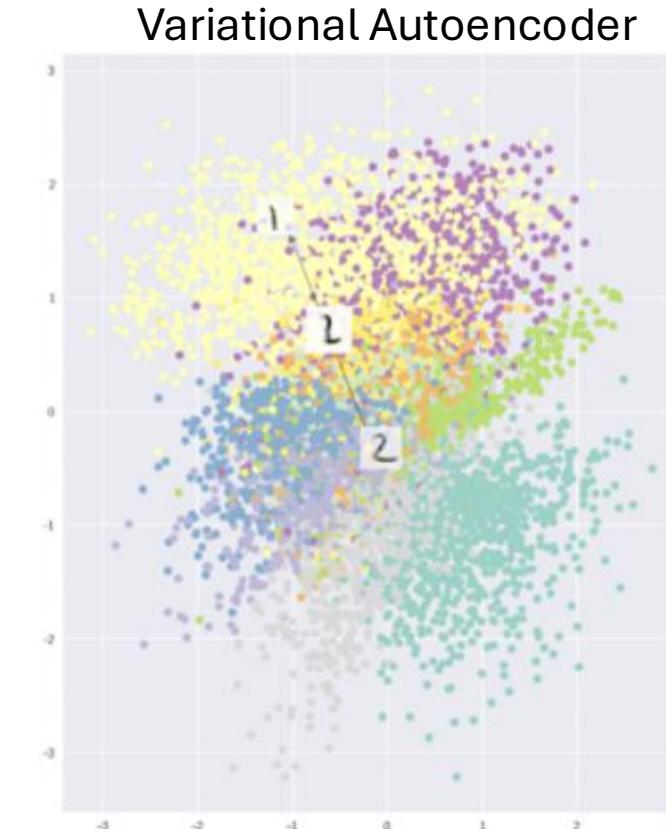
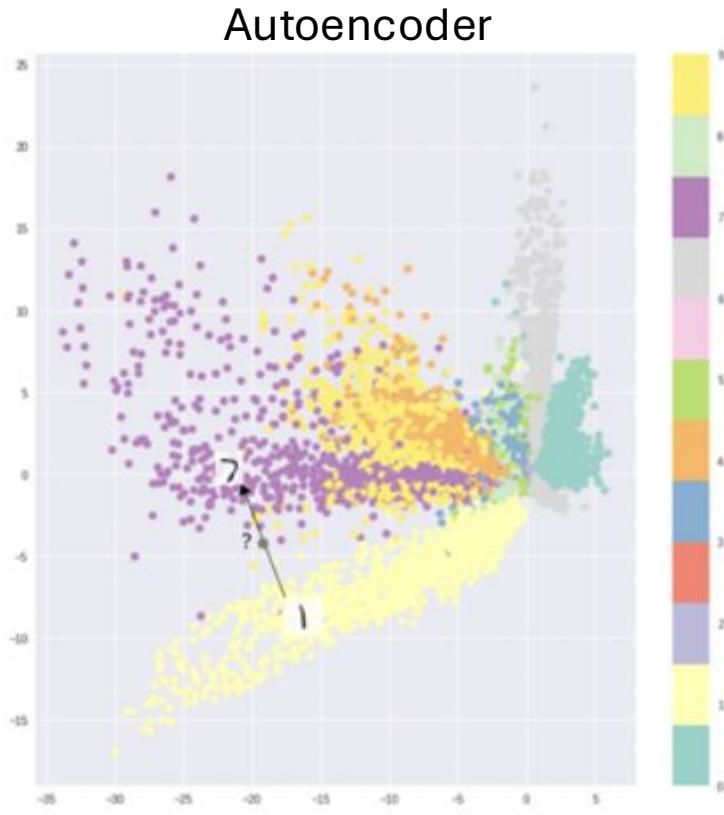
How to address issues with overfitting outputs? Try to learn more *variation* in outputs.



Explained variance: PC1=0.32, PC2=0.11

# Issues with Autoencoders

What might a better latent space look like for generation?



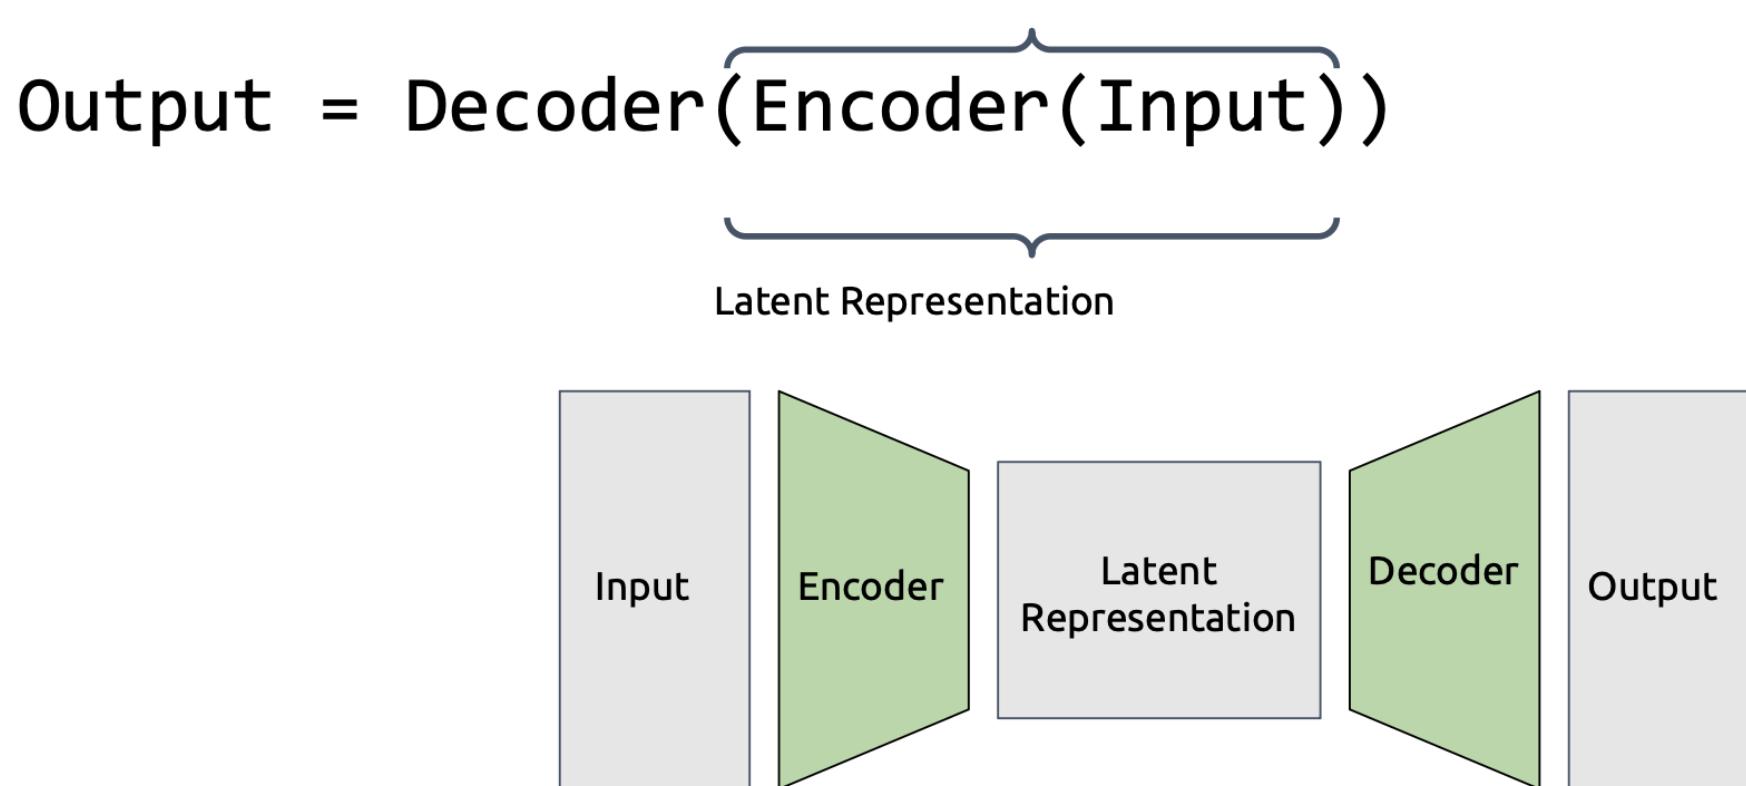
# Variational Autoencoders

Autoencoder's goal: Reconstruct the original input

Variational Autoencoder's goal: Generate a new output that resembles the input

# Building up the VAE Architecture

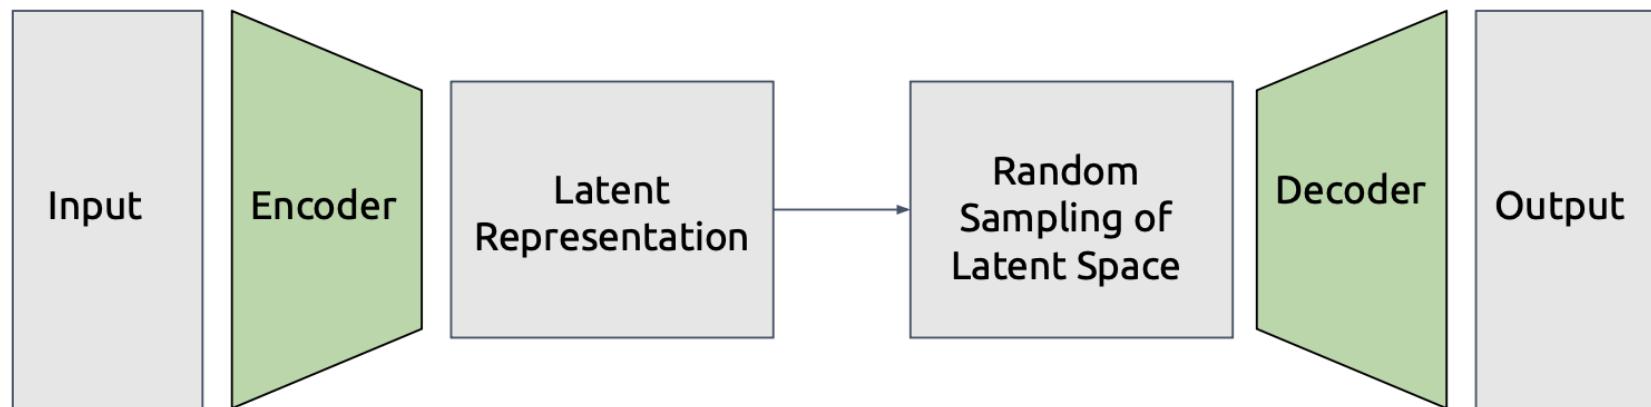
If we were to describe an autoencoder functionally:



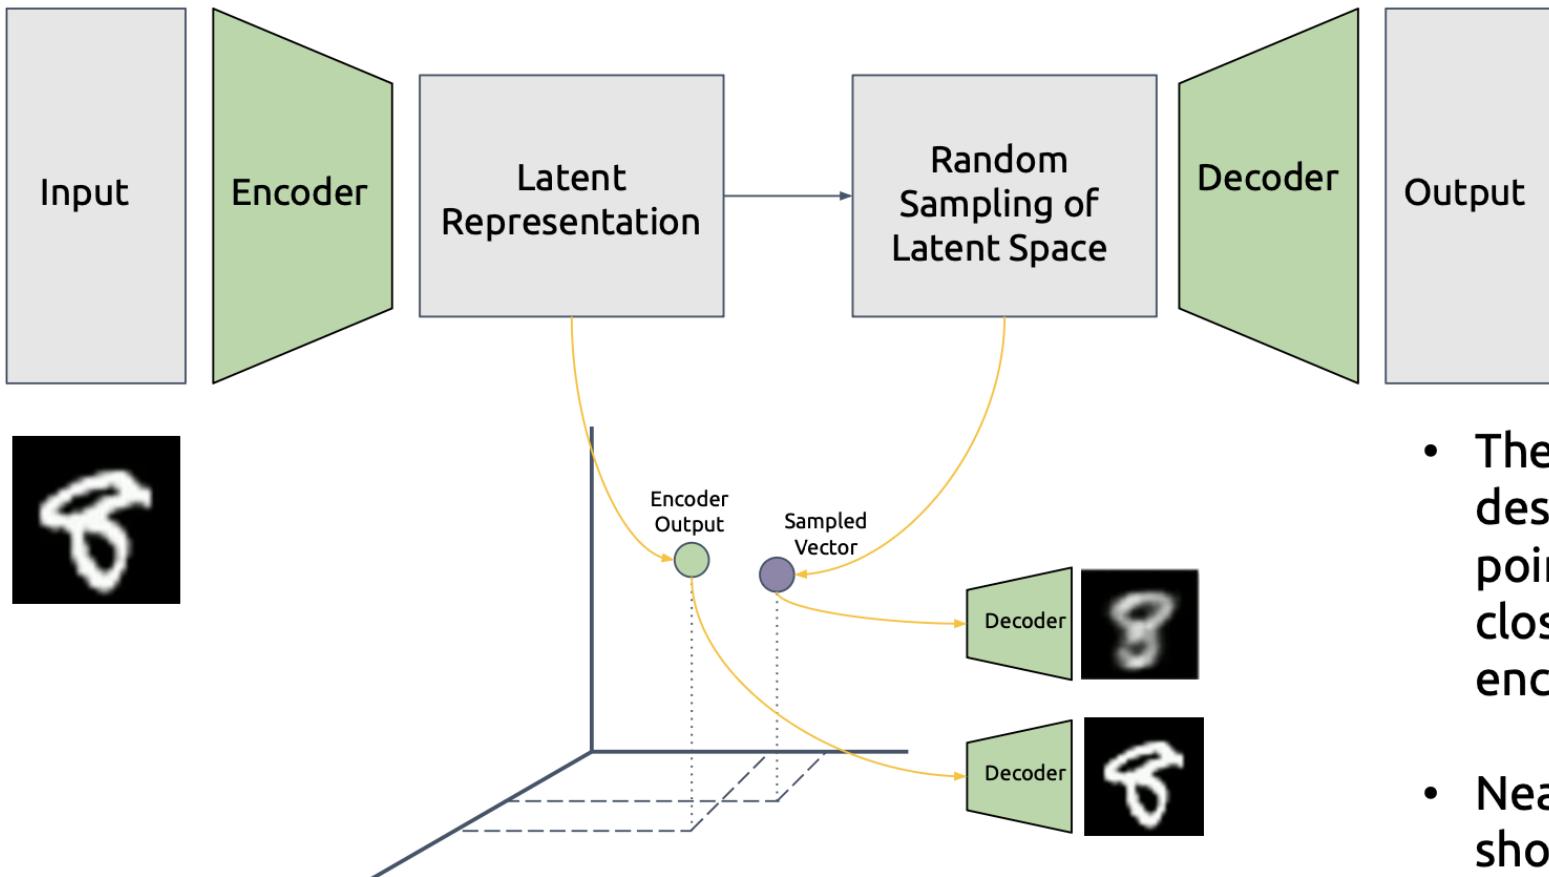
# Building up the VAE Architecture

For variational autoencoders, we also do a random sampling operation at the bottleneck

```
Output = Decoder(random_sample(Encoder(Input)))
```



# How does random sampling in latent space lead to variation?



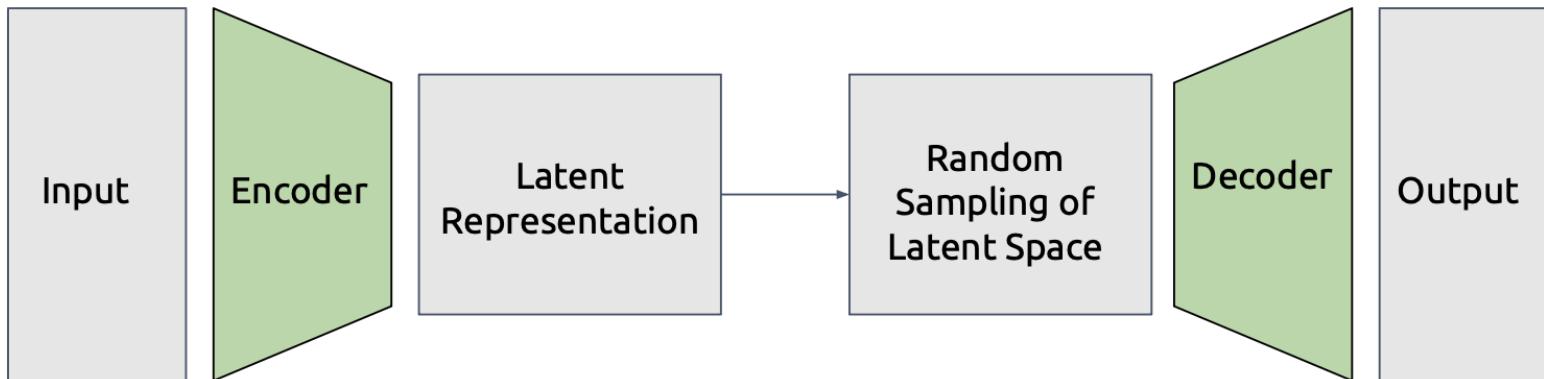
- The random sampling should be designed to produce random points in latent space that are close to the output of the encoder
- Nearby points in the latent space should decode to similar images

# How should `random_sample` be defined?

```
Output = Decoder(random_sample(Encoder(Input)))
```

- We want the sample to be close to the encoder output
- One option: sample from a Gaussian centered at Encoder(Input)

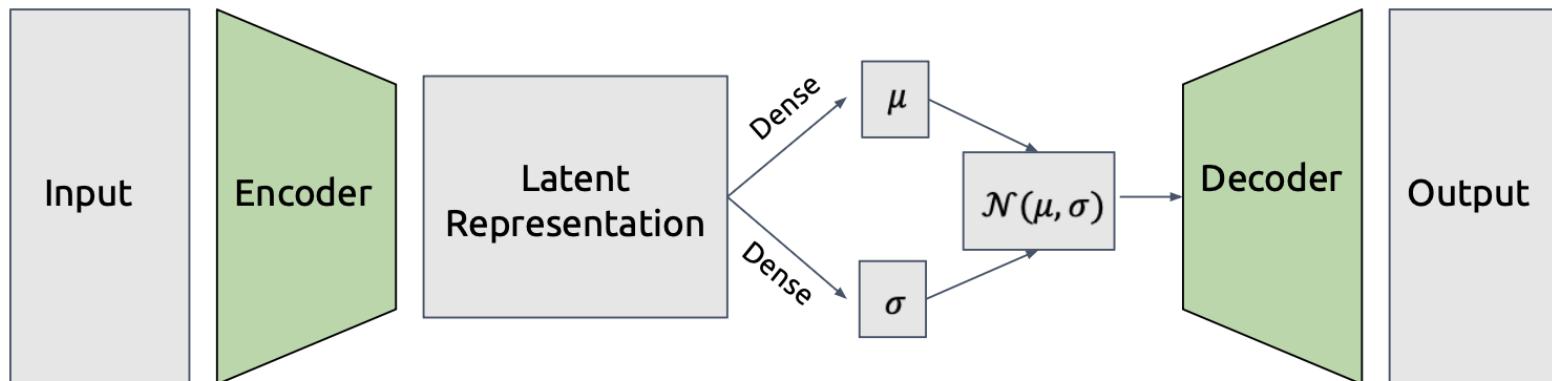
What can we modify?



# How should `random_sample` be defined?

```
Output = Decoder(random_sample(Encoder(Input)))
```

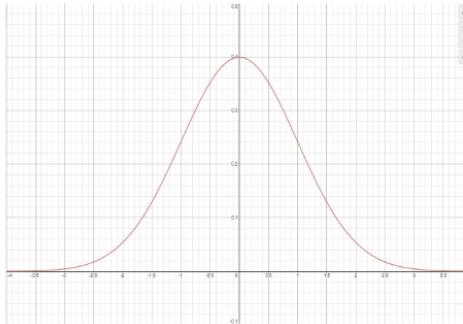
- We want the sample to be close to the encoder output
- One option: sample from a Gaussian centered at `Encoder(Input)`
- Use two dense layers to convert the encoder output into the mean and standard deviation of the Gaussian



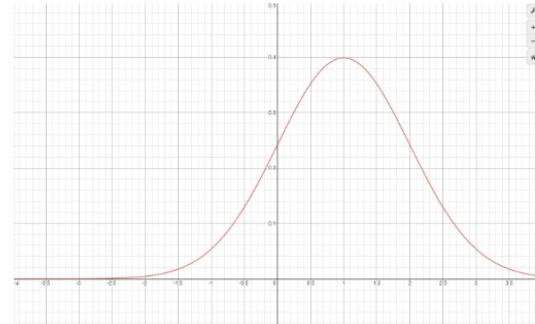
Any questions?



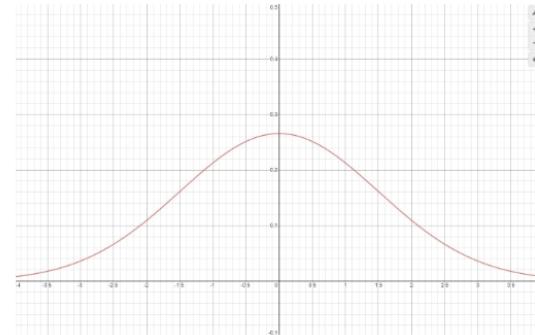
# How should `random_sample` be defined?



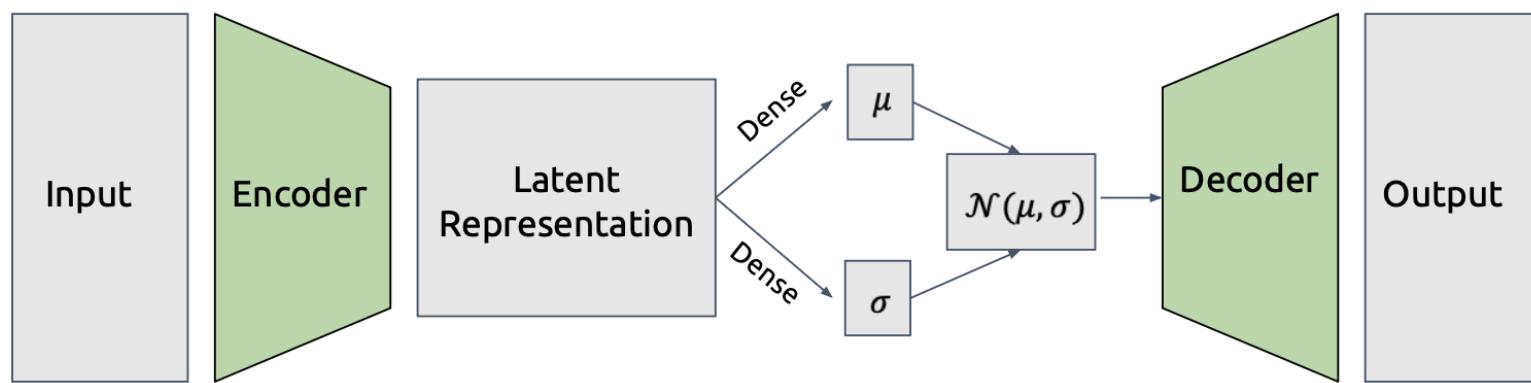
$$\begin{aligned}\mu &= 0 \\ \sigma &= 1\end{aligned}$$



$$\begin{aligned}\mu &= 1 \\ \sigma &= 1\end{aligned}$$



$$\begin{aligned}\mu &= 0 \\ \sigma &= 1.5\end{aligned}$$

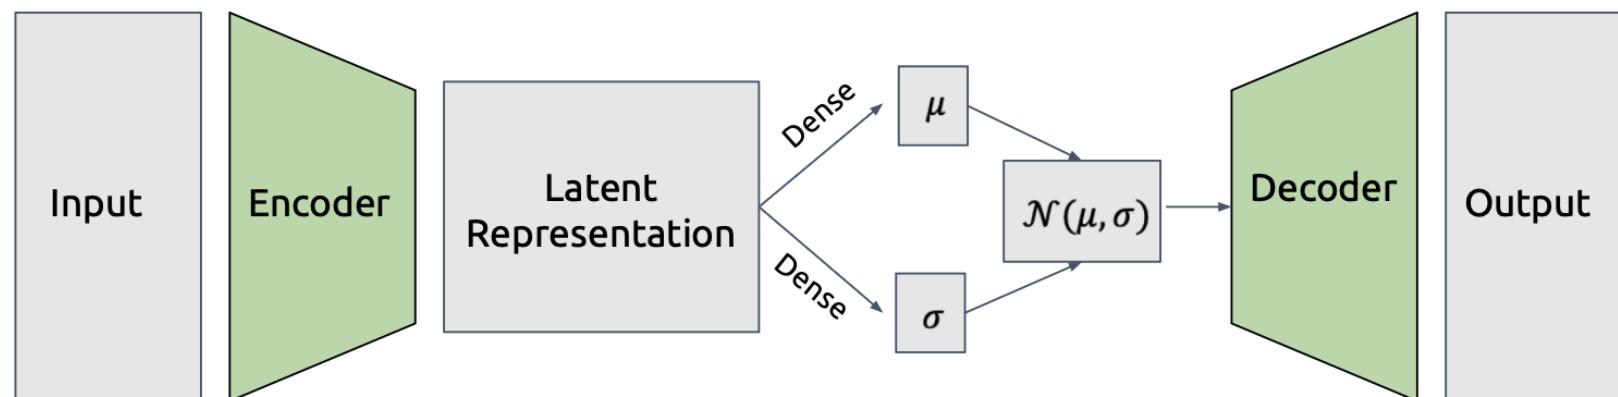


# Training a VAE

Two goals:

1. Reproduce an output similar to the input ( $\text{Input} \approx \text{Output}$ )
2. Have some variation in our output ( $\text{Input} \neq \text{Output}$ )

- Seems like two conflicting goals!
- How do we resolve these two goals?



# Weighted Combination of Losses

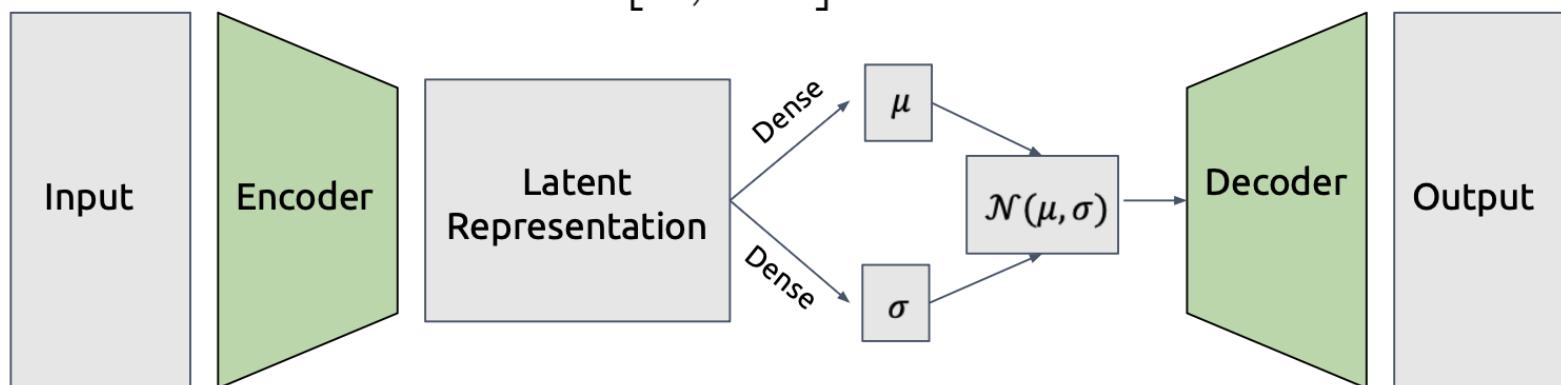
$L_1$  = loss associated with producing output similar to input

$L_2$  = loss associated with producing output with some variation to input

$$L = L_1 + \lambda L_2$$

Total Loss:

$$\lambda \in [0, \infty]$$



# VAE Losses

$L_1$  is easy, we've seen this before

$$L_1(x, \hat{x}) = \|x - \hat{x}\|_2 \text{ (MSE)}$$

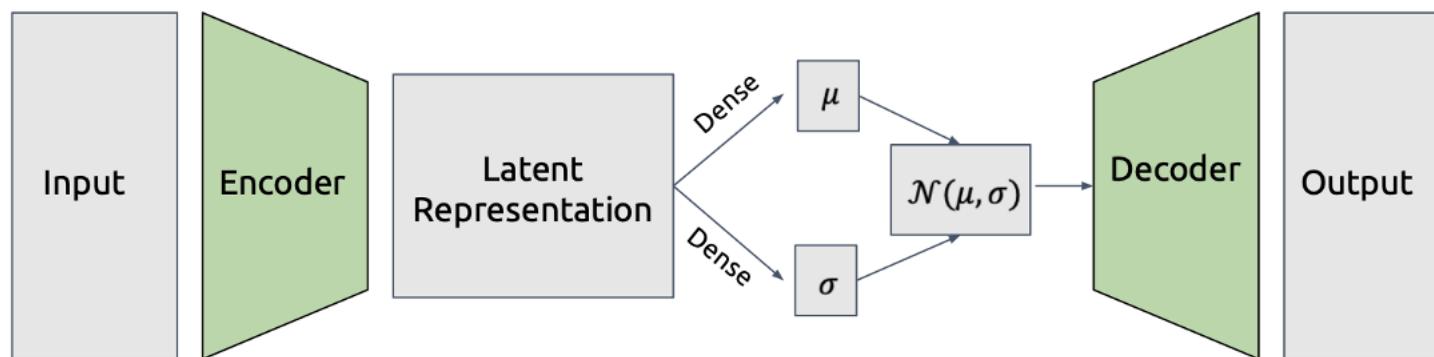
But what is  $L_2$ ? How do we measure how much variation our output has?

$$L_2(??, ??) = ???$$

# Defining the Variation Loss

Whatever our loss function, it needs to encourage  $\sigma > 0$ , or else the model will force  $\sigma$  to 0 in an effort to create the best recreations possible.

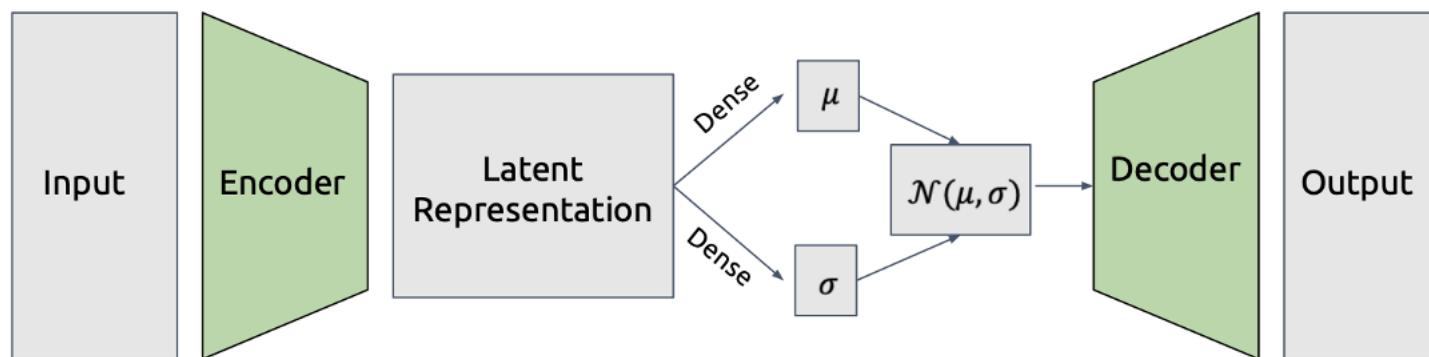
If  $\sigma = 0$ , then the VAE will behave the same as an autoencoder!



# Defining the Variation Loss

Whatever our loss function, it needs to encourage  $\sigma > 0$ , or else the model will force  $\sigma$  to 0 in an effort to create the best recreations possible.

But it can't be too big... because too much variation will create poor reconstructions.

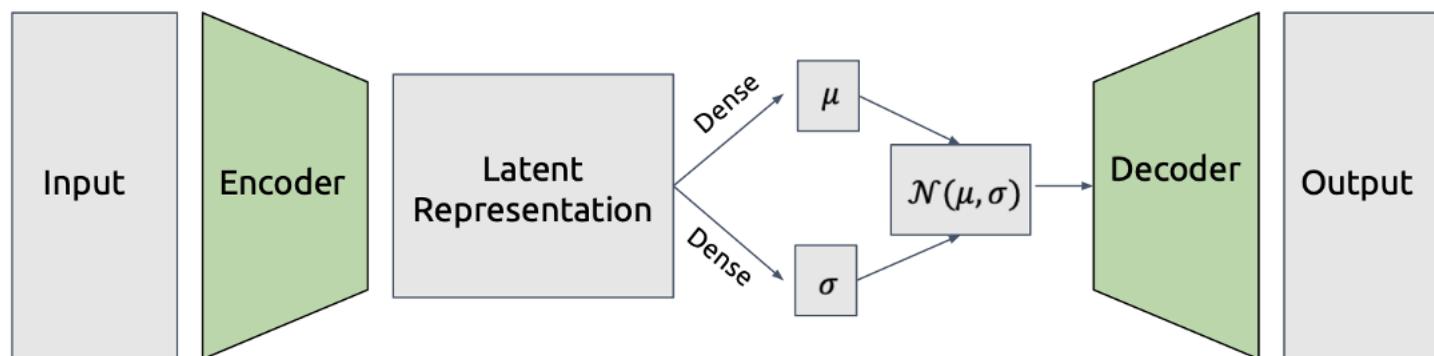


# Defining the Variation Loss

Whatever our loss function, it needs to encourage  $\sigma > 0$ , or else the model will force  $\sigma$  to 0 in an effort to create the best recreations possible.

But it can't be too big... because too much variation will create poor reconstructions.

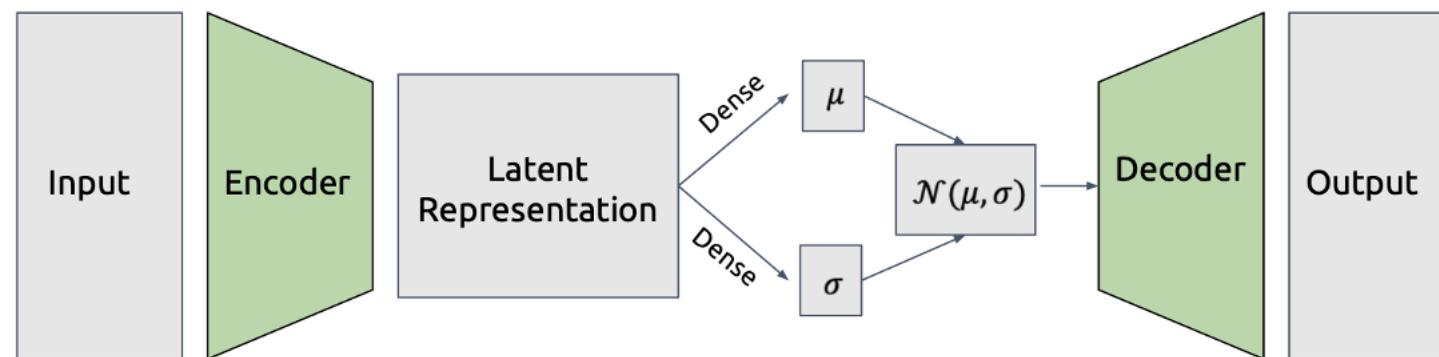
Also, what should  $\mu$  be?



# Defining the Variation Loss

The idea:

- Introduce a *prior* probability function we want our latent space to look like.
- Encourage  $N(\mu, \sigma)$  close to  $N(0, 1)$
- (This will have beneficial properties we'll see later)



# How do we measure distance between probabilities?

## Kullback–Leibler (KL) Divergence

$$D_{KL}(P||Q) = \int_{-\infty}^{\infty} p(x) \log \left( \frac{p(x)}{q(x)} \right) dx$$

What this says:

- “**Everywhere that  $p$  has probability density...**”
- “...**the difference between  $p$  and  $q$  should be small**”
  - Difference in log probabilities (remember that  $\log\left(\frac{a}{b}\right) = \log(a) - \log(b)$ )

# KL Divergence

- Expensive to compute, in general (no closed form, have to numerically approximate the integral)
- But! There is a closed form for Gaussians:

$$D_{KL}(\mathcal{N}(\mu, \sigma^2) || \mathcal{N}(0, 1)) = \frac{1}{2} \sum_{i=1}^k (\mu_i^2 + \sigma_i^2 - \ln \sigma_i^2 - 1)$$

K is the dimensionality of  $\vec{\mu}, \vec{\sigma}$  (i.e., the size of the encoding)

# The Final VAE Loss Function

We now have all the tools necessary to construct our loss function.

$$L = L_1 + \lambda L_2 \quad \lambda \in [0, \infty]$$

Which turns into this:

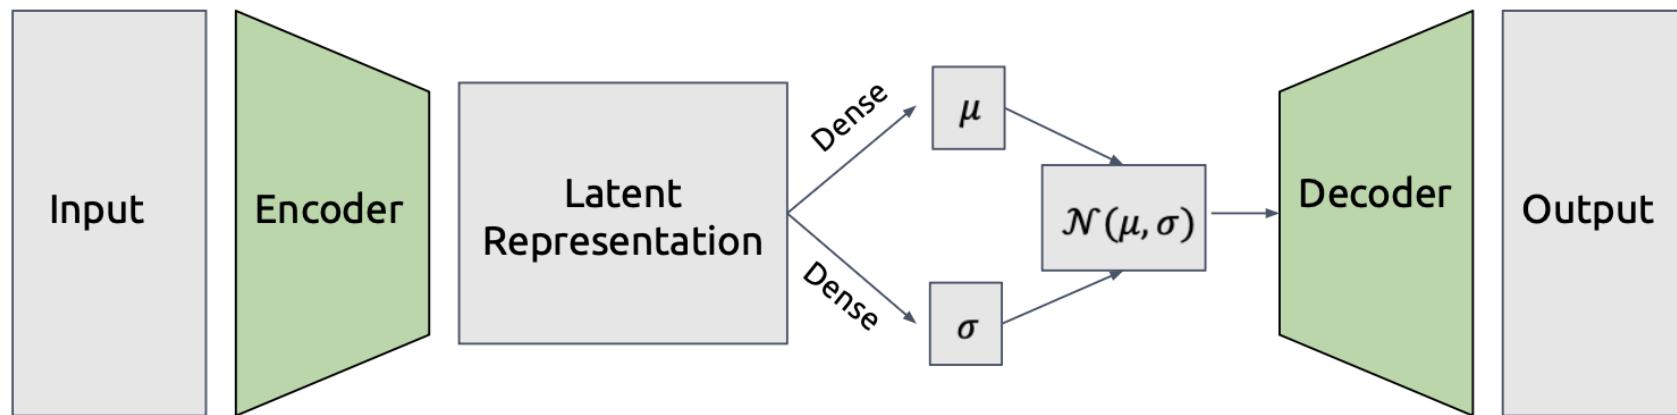
$$L = ||x - \hat{x}||_2^2 + \lambda D_{KL}(\mathcal{N}(\mu, \sigma), \mathcal{N}(0, 1))$$

Any questions?



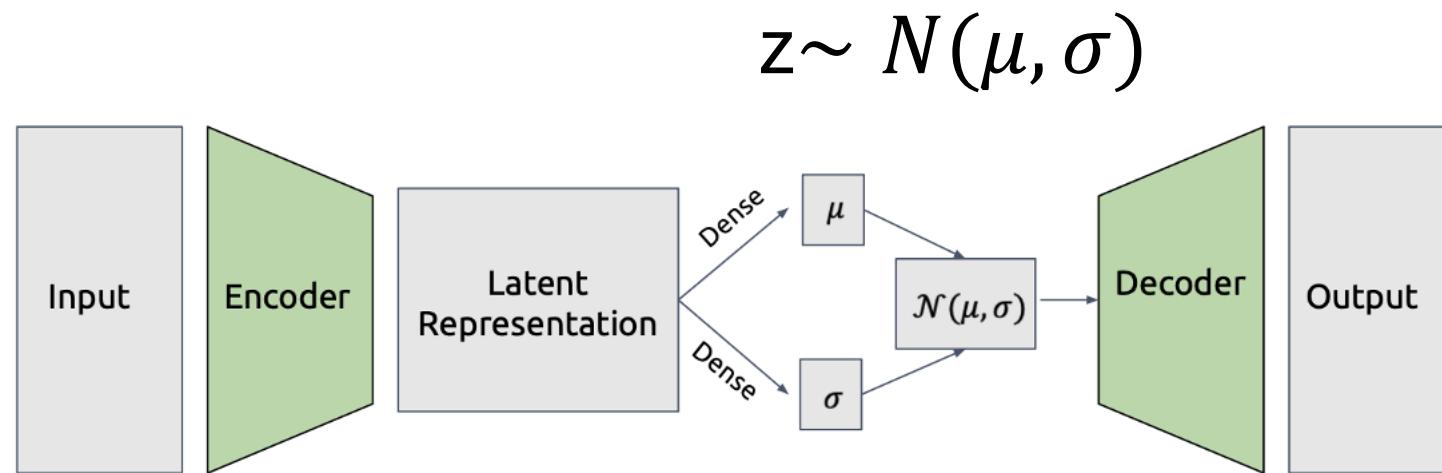
# Putting it all together

$$L = \|x - \hat{x}\|_2^2 + \lambda D_{KL}(\mathcal{N}(\mu, \sigma), \mathcal{N}(0, 1))$$



# There's just one issue

How do we take the gradient of a sampling operation?



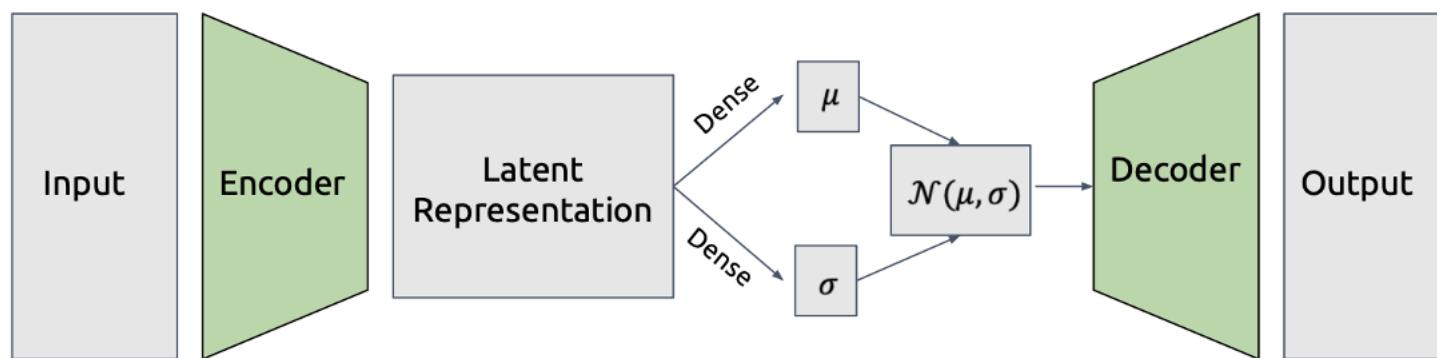
# Reparametrization Trick

$$z \sim N(\mu, \sigma)$$

Can be rewritten as:

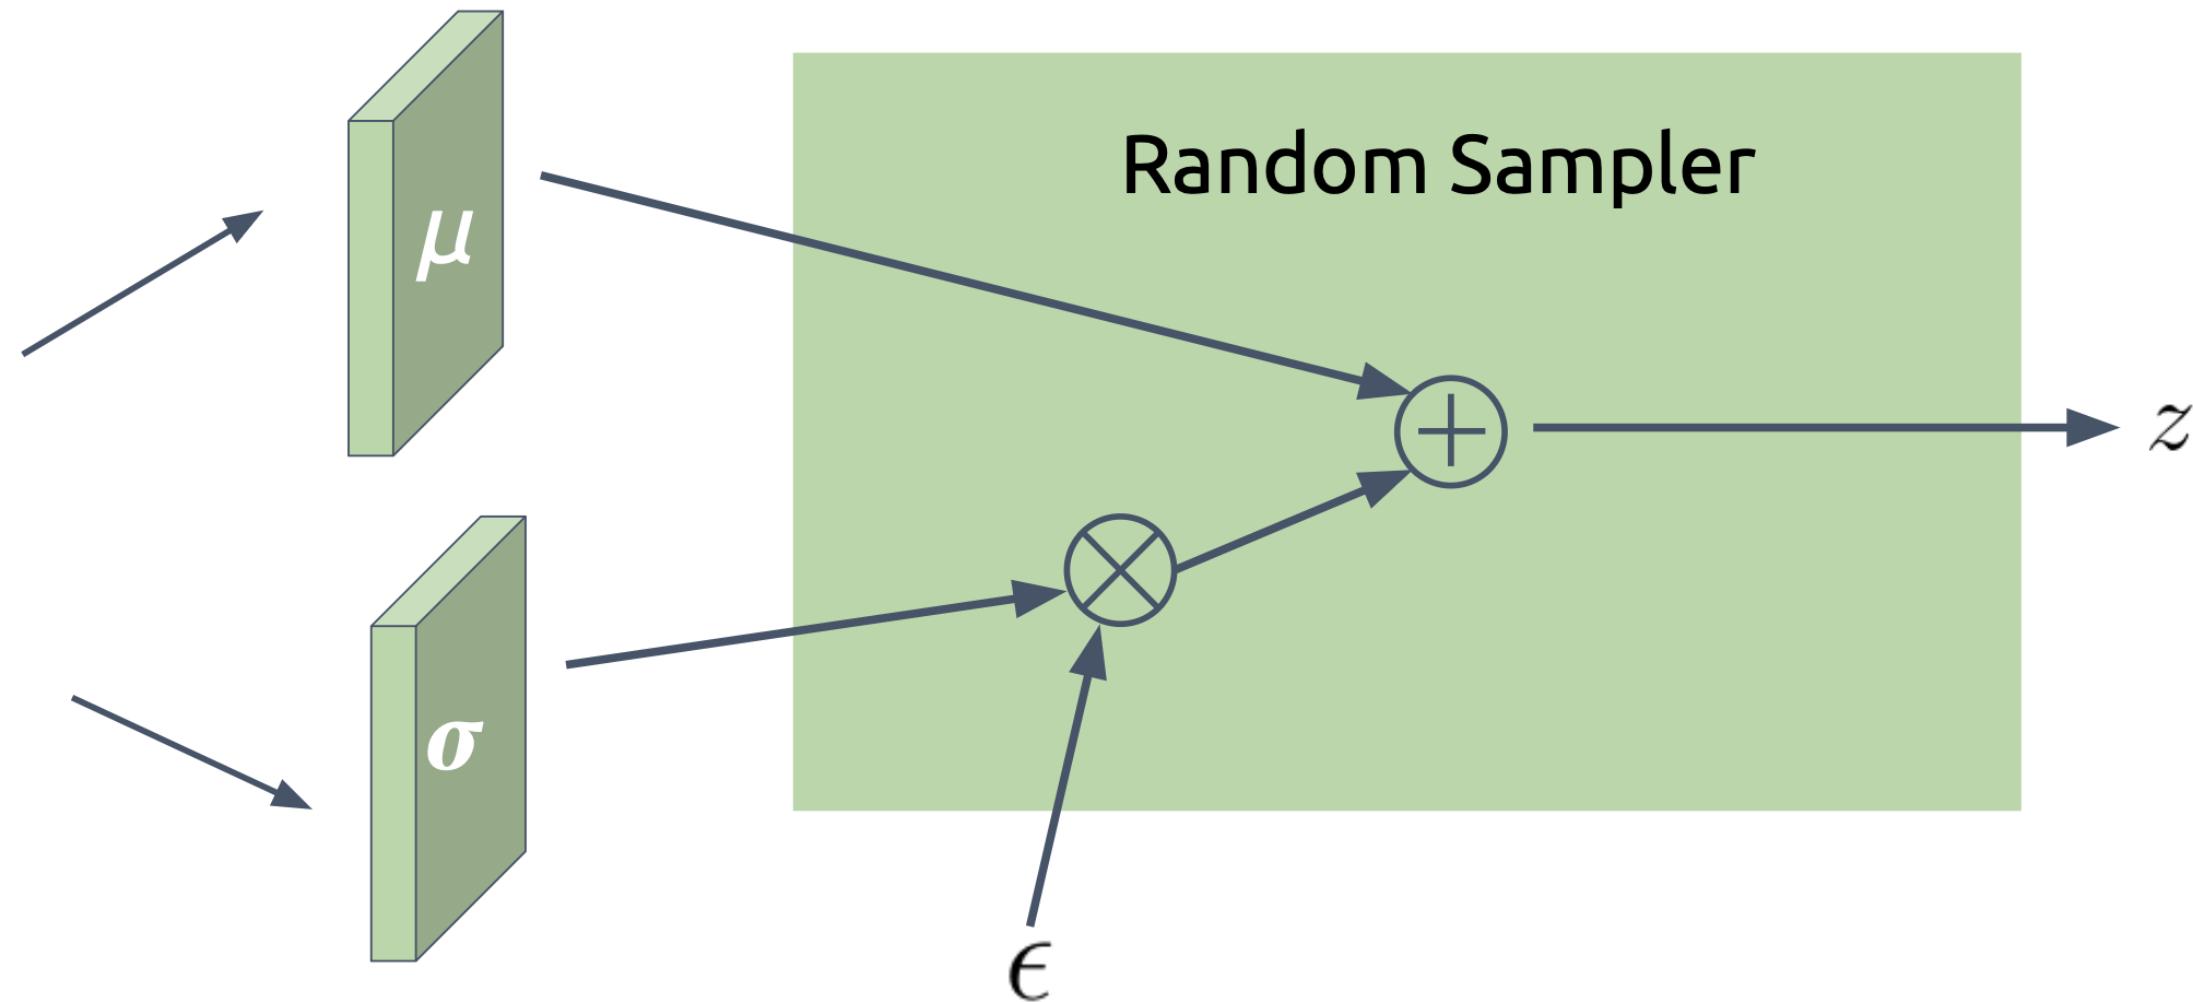
$$z = \mu + \epsilon\sigma, \text{ where } \epsilon \sim N(0, 1)$$

Random sampling operation ( $\epsilon$ ) no longer depends on learnable parameters

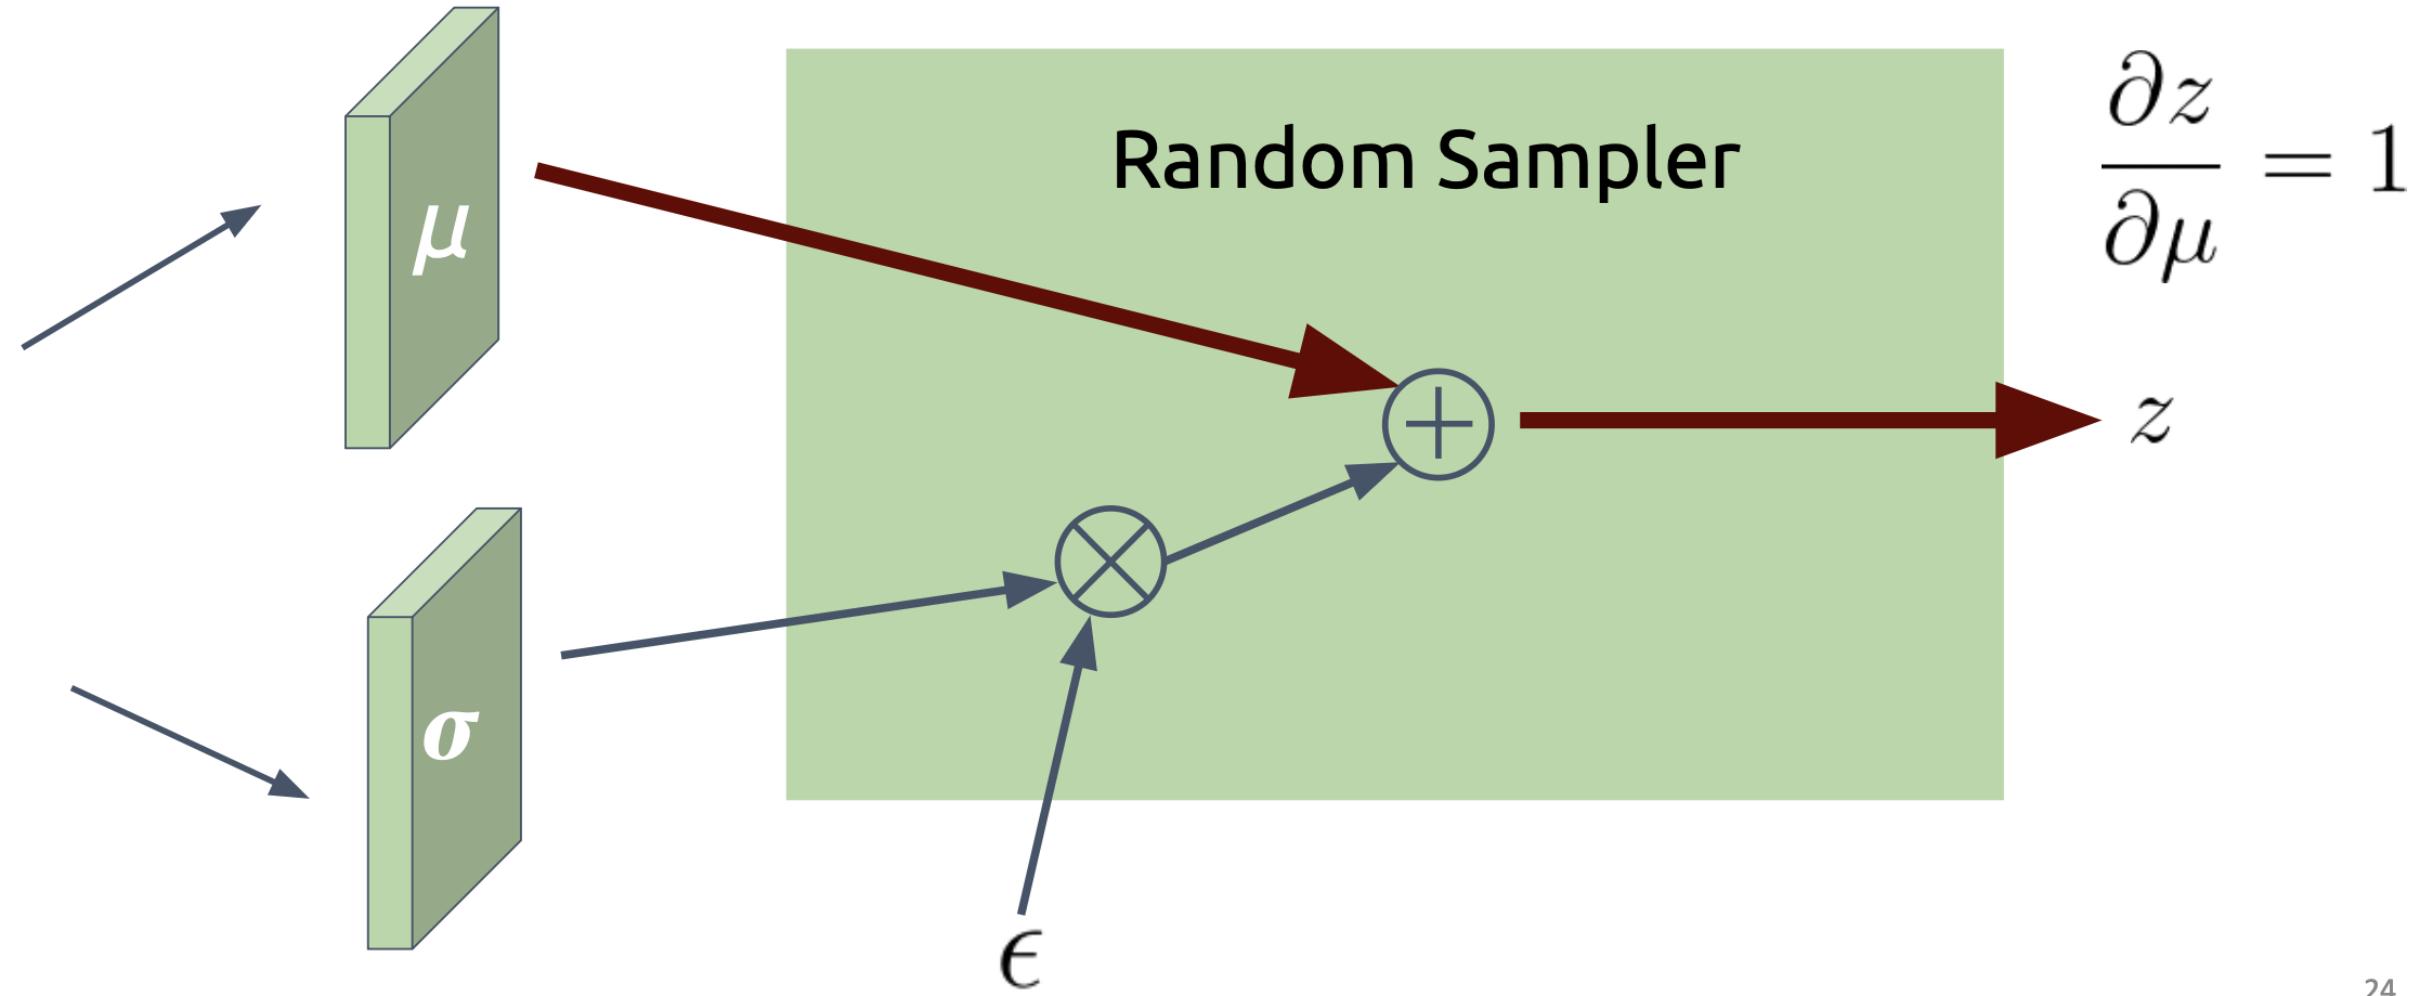


Another explanation of why this is needed: <https://gregorygundersen.com/blog/2018/04/29/reparameterization/>

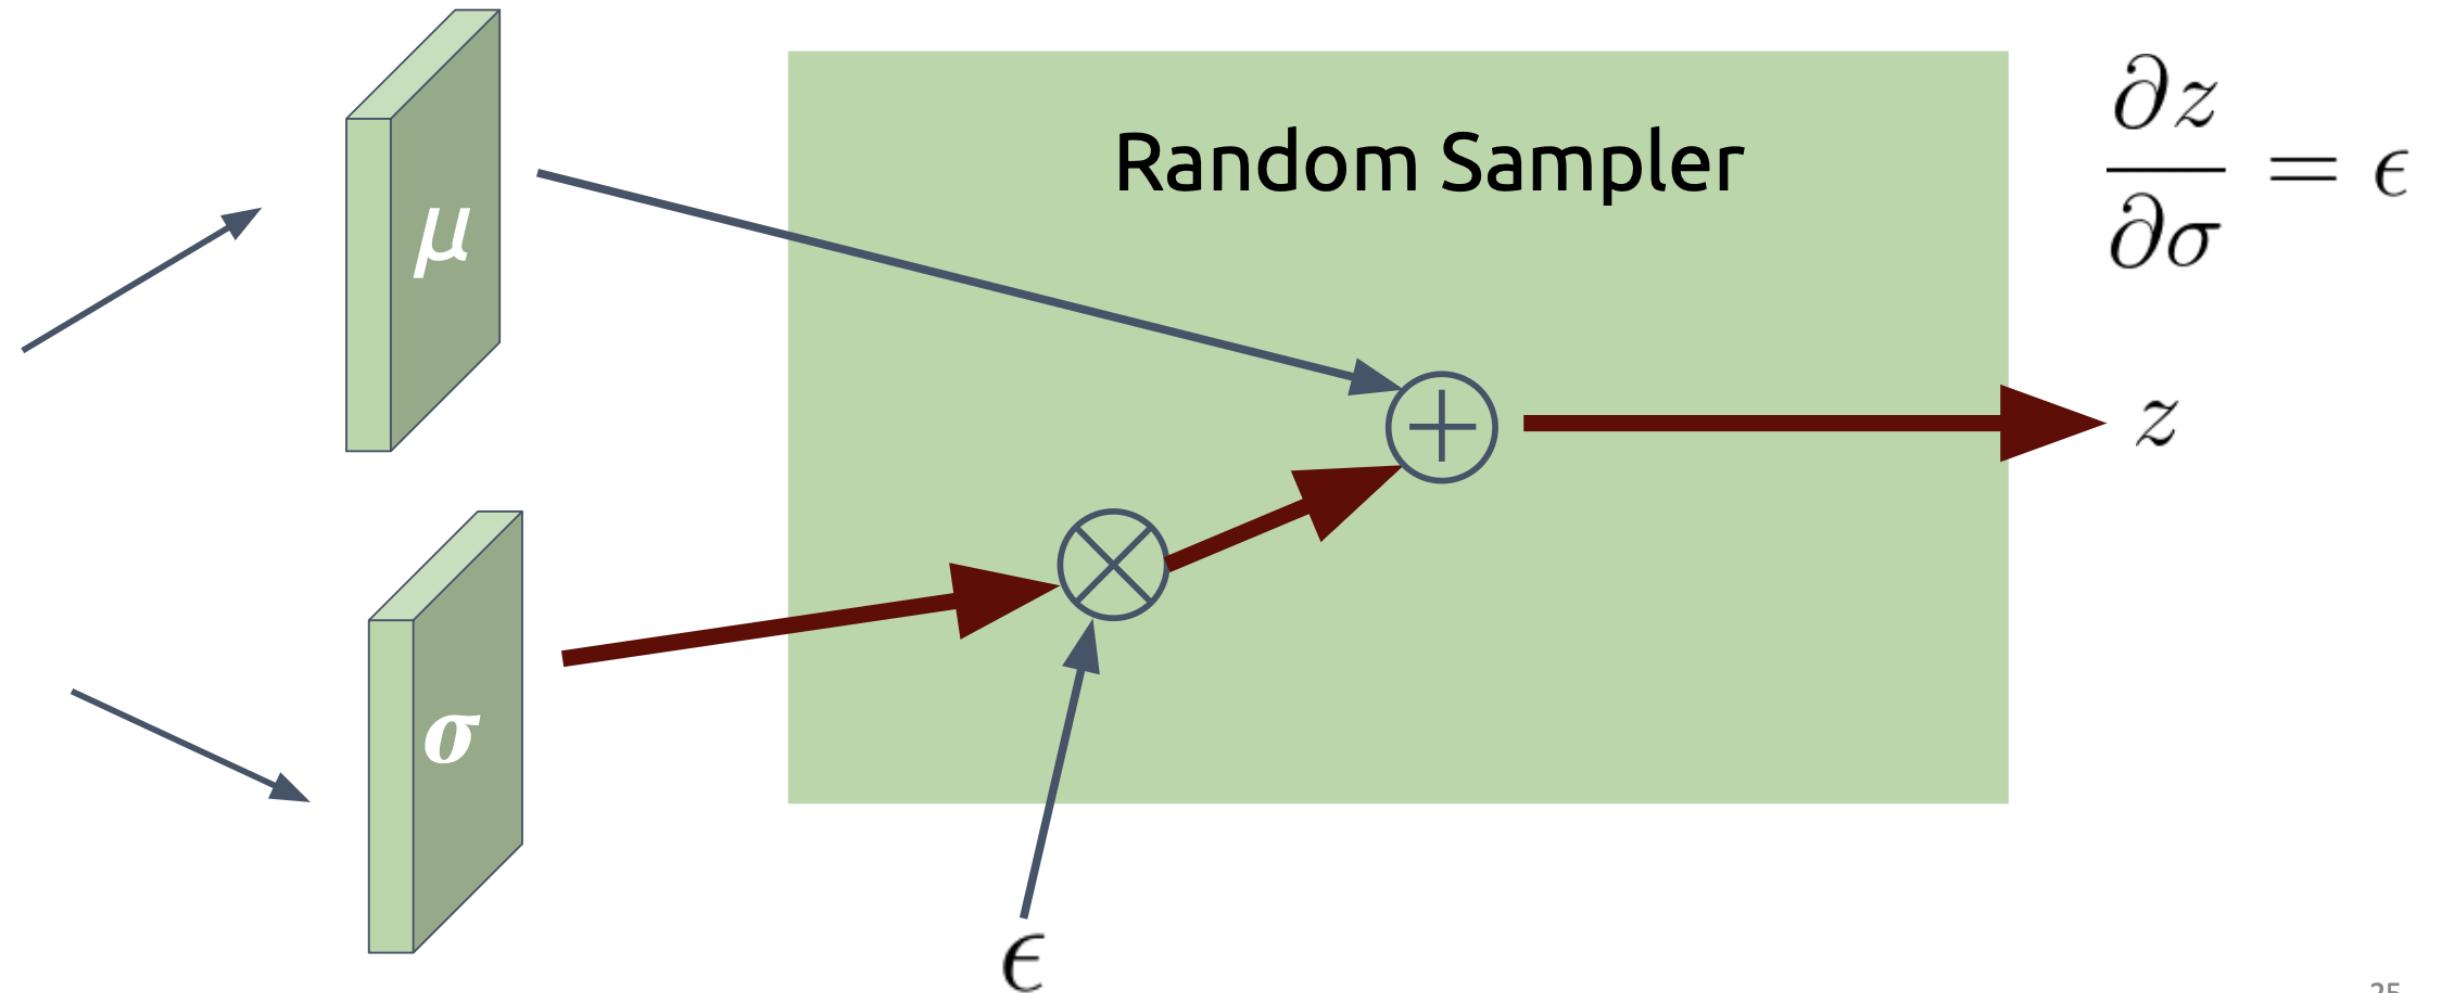
# Random Sampler with Reparameterization Trick



# Random Sampler with Reparameterization Trick

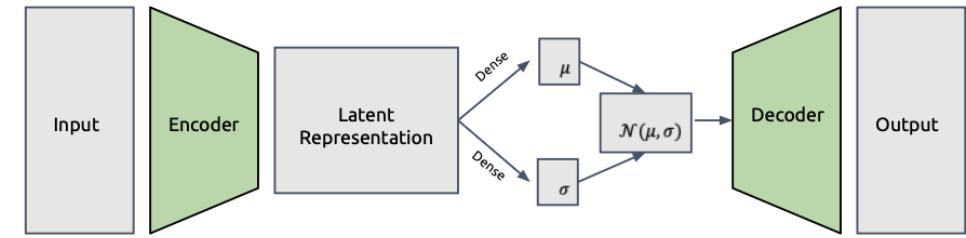


# Random Sampler with Reparameterization Trick



# One more practical detail

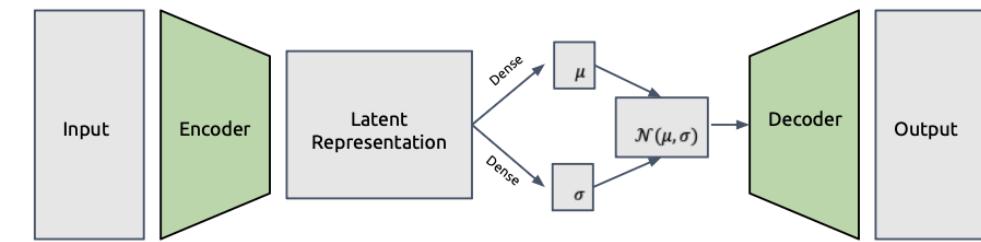
Let's again consider our sampling operation



$$z \sim \mathcal{N}(\mu, \sigma)$$
$$\mu_i \in [-\infty, \infty] \quad \sigma_i \in [0, \infty]$$

- Nothing prevents the neural network from outputting ***negative*** values for the standard deviation.
- Instead of predicting  $\sigma$ , we will instead predict  $\log(\sigma^2)$ . This ensures that every  $\sigma_i \in [0, \infty]$ 
  - i.e. just treat the output of the Dense layer as if it is  $\log(\sigma^2)$

# One more practical detail



Let's again consider our sampling operation

$$z \sim \mathcal{N}(\mu, \sigma)$$
$$\mu_i \in [-\infty, \infty] \quad \sigma_i \in [0, \infty]$$

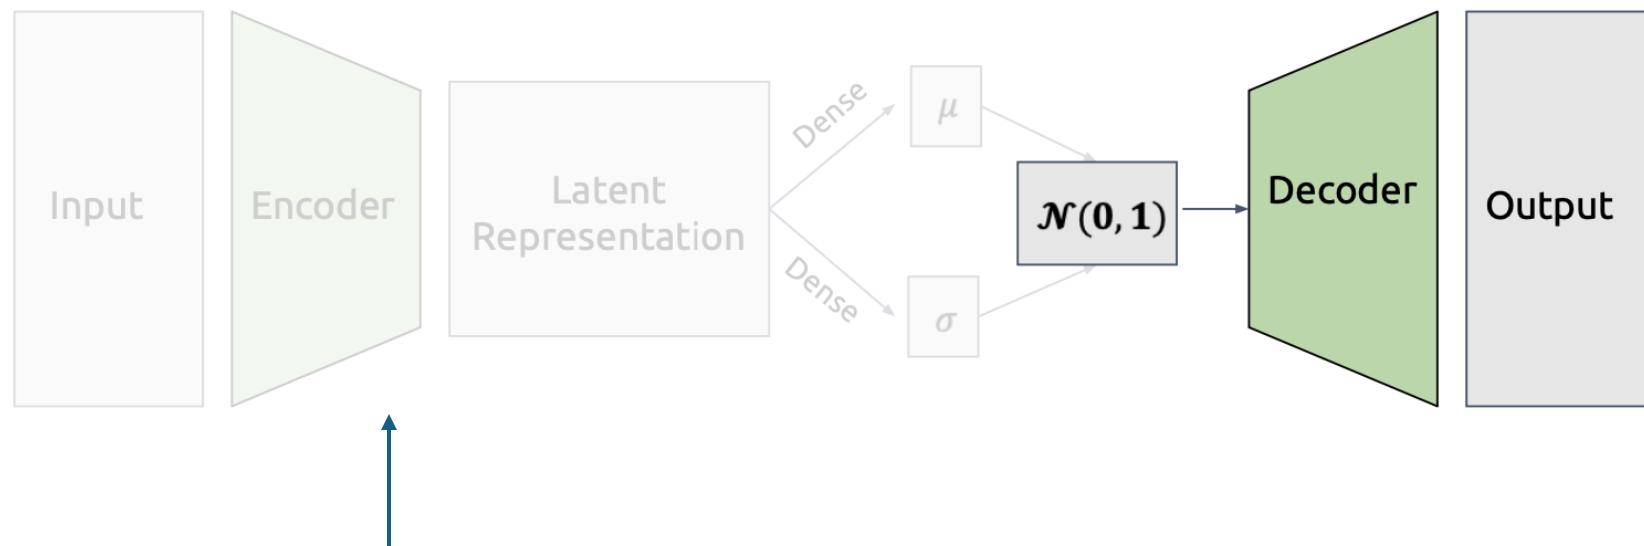
Any questions?



- Instead of predicting  $\sigma$ , we will instead predict  $\log(\sigma^2)$ . This ensures that every  $\sigma_i \in [0, \infty]$ 
  - i.e. just treat the output of the Dense layer as if it is  $\log(\sigma^2)$

$$D_{KL}(\mathcal{N}(\mu, \sigma^2) || \mathcal{N}(0, 1)) = \frac{1}{2} \sum_{i=1}^k (\mu_i^2 + \sigma_i^2 - \ln \sigma_i^2 - 1)$$

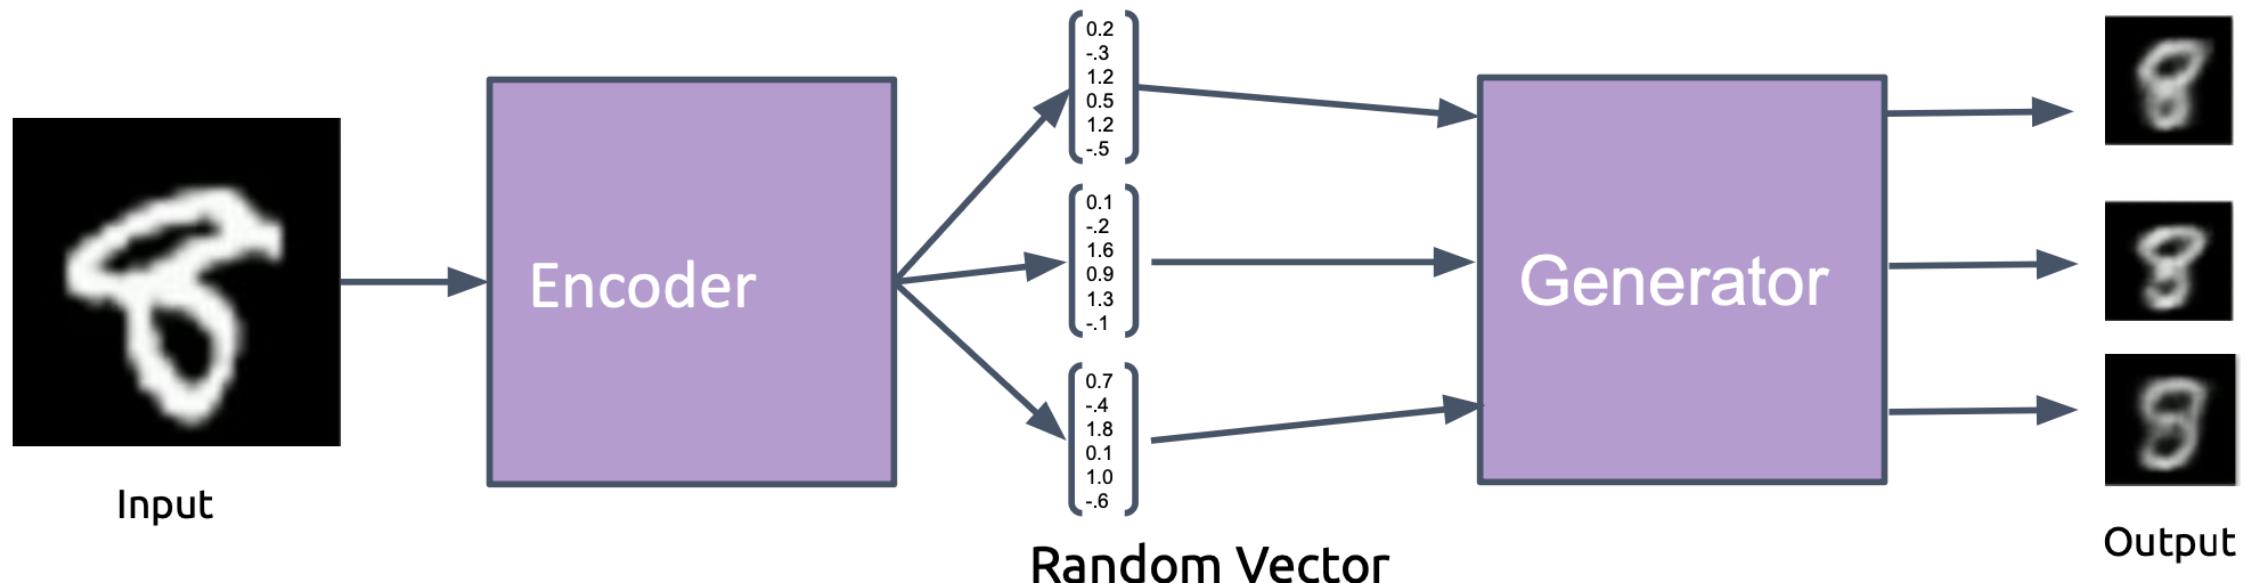
# Sampling from a VAE



- Discard this part of the network...
- ...and set  $(\mu, \sigma) = (0, 1)$

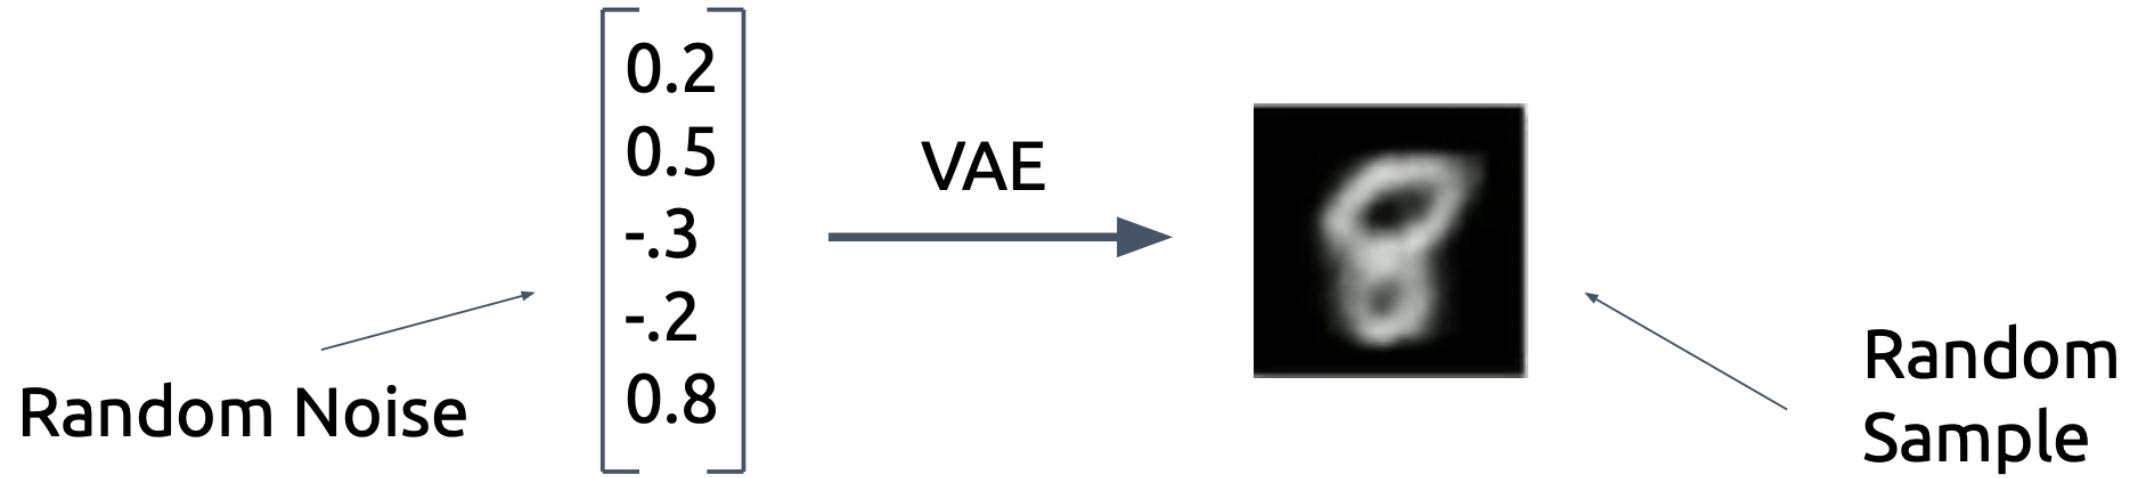
# Sampling from a VAE

- We can use a trained VAE to generate random variants of an input data point...



# Sampling from a VAE

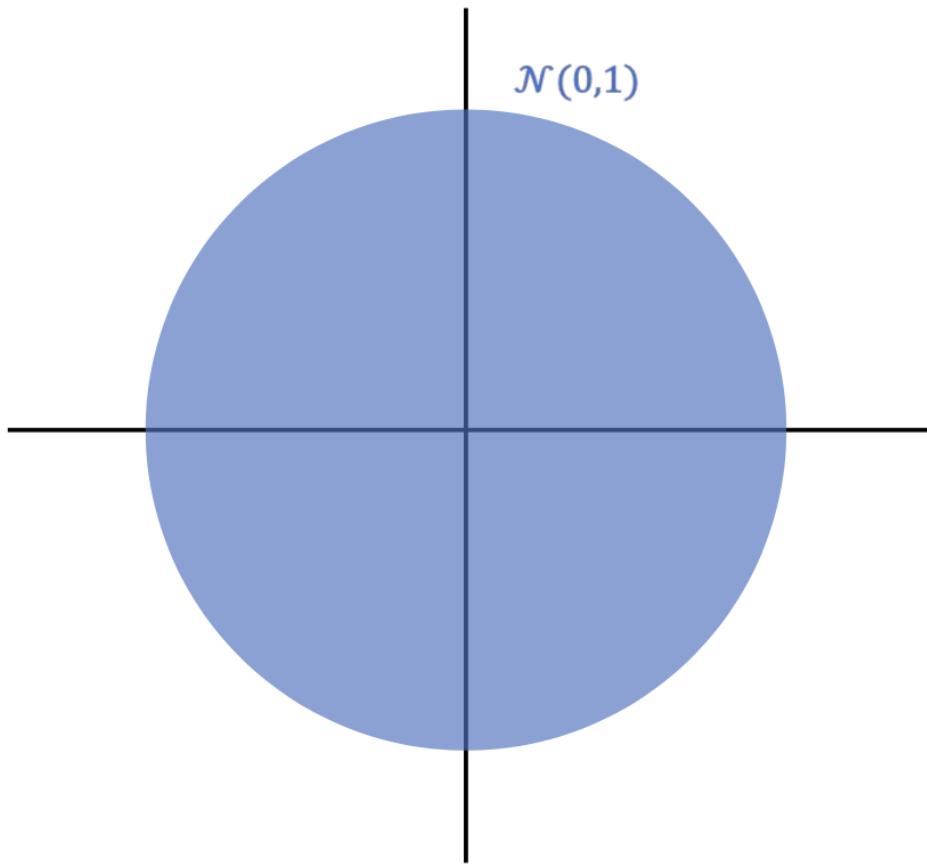
... But ultimately, we want to draw random samples from a VAE



***How can we do this?***

This is where our particular choice of training loss will pay off

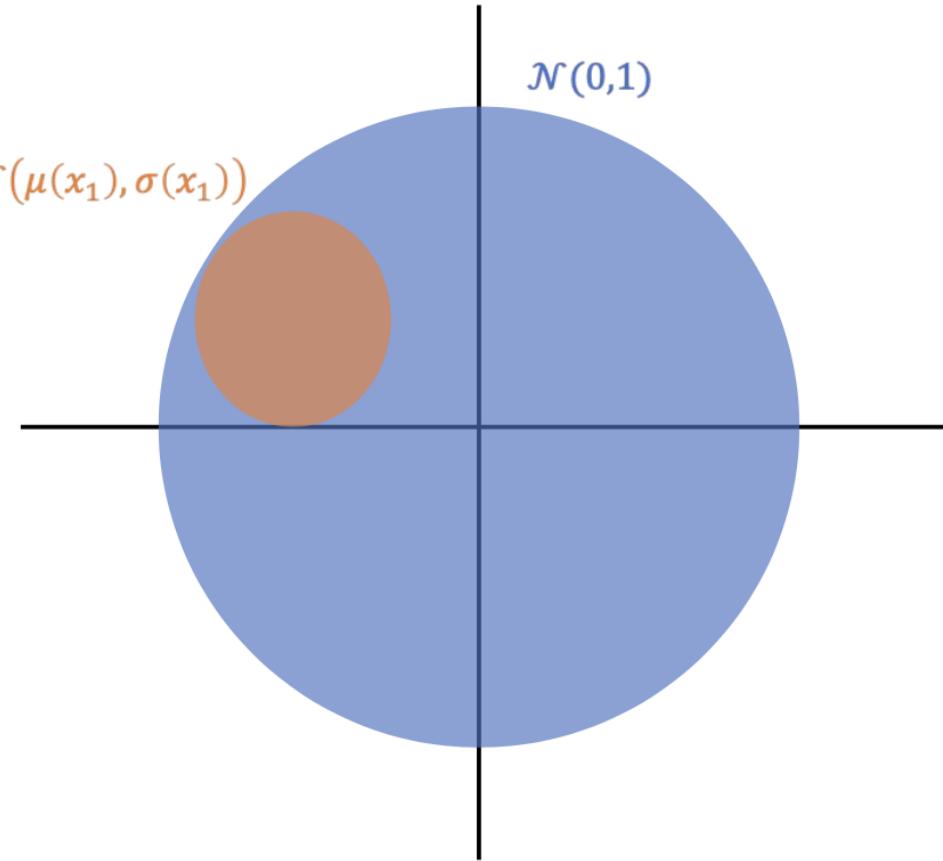
# Encoding different points into latent space



Let this circle represent the probability density of a unit Gaussian in latent space

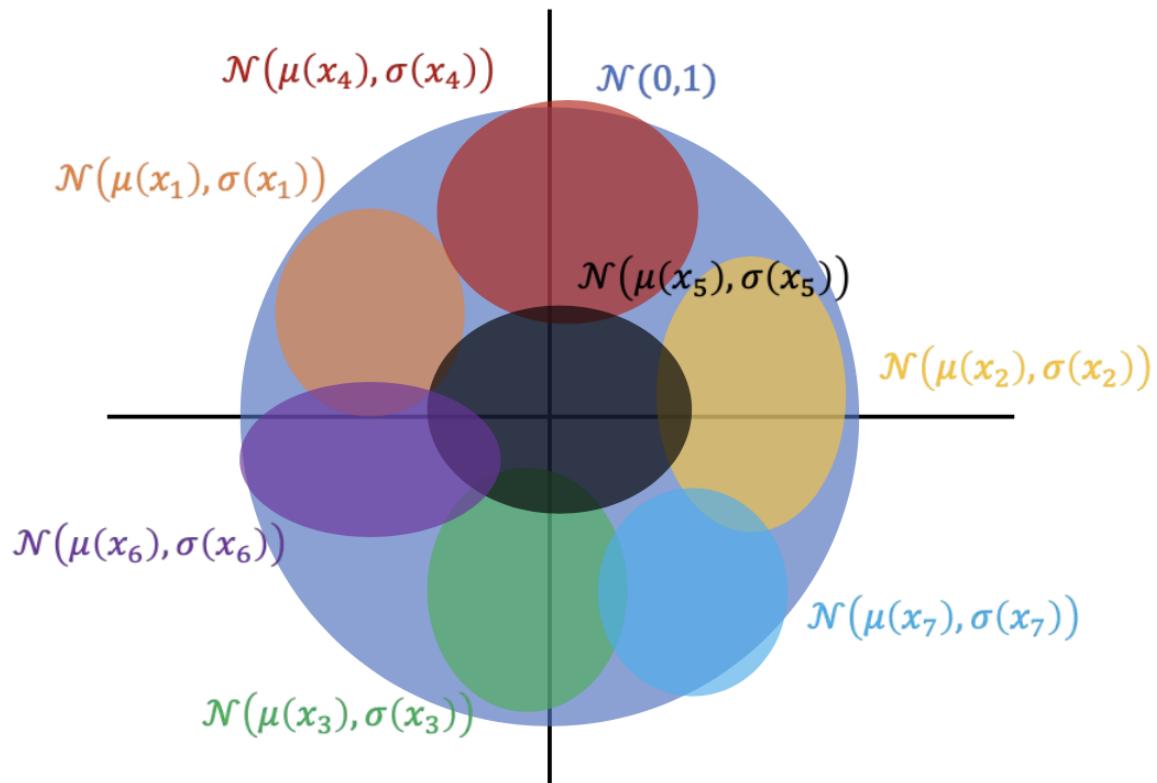
# Encoding different points into latent space

Let this circle represent the probability density of the  $\mathcal{N}(\mu, \sigma)$  distribution that the encoder predicts given an input data point  $x_1$



# Encoding different points into latent space

$$L = \|x - \hat{x}\|_2^2 + \lambda D_{KL}(\mathcal{N}(\mu, \sigma), \mathcal{N}(0, 1))$$

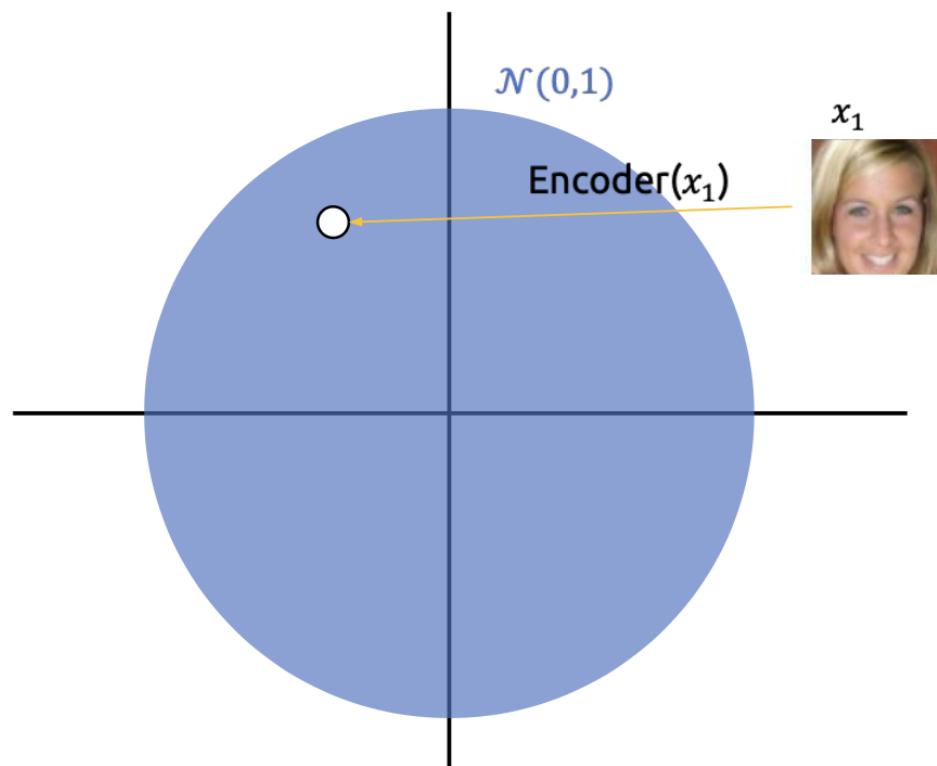


Because of our KL divergence loss, the  $\mathcal{N}(\mu, \sigma)$  for any input data point has to be somewhat similar to  $\mathcal{N}(0,1)$

So, if we sample a point from  $\mathcal{N}(0,1)$ , it is very likely to fall within one of these encoded

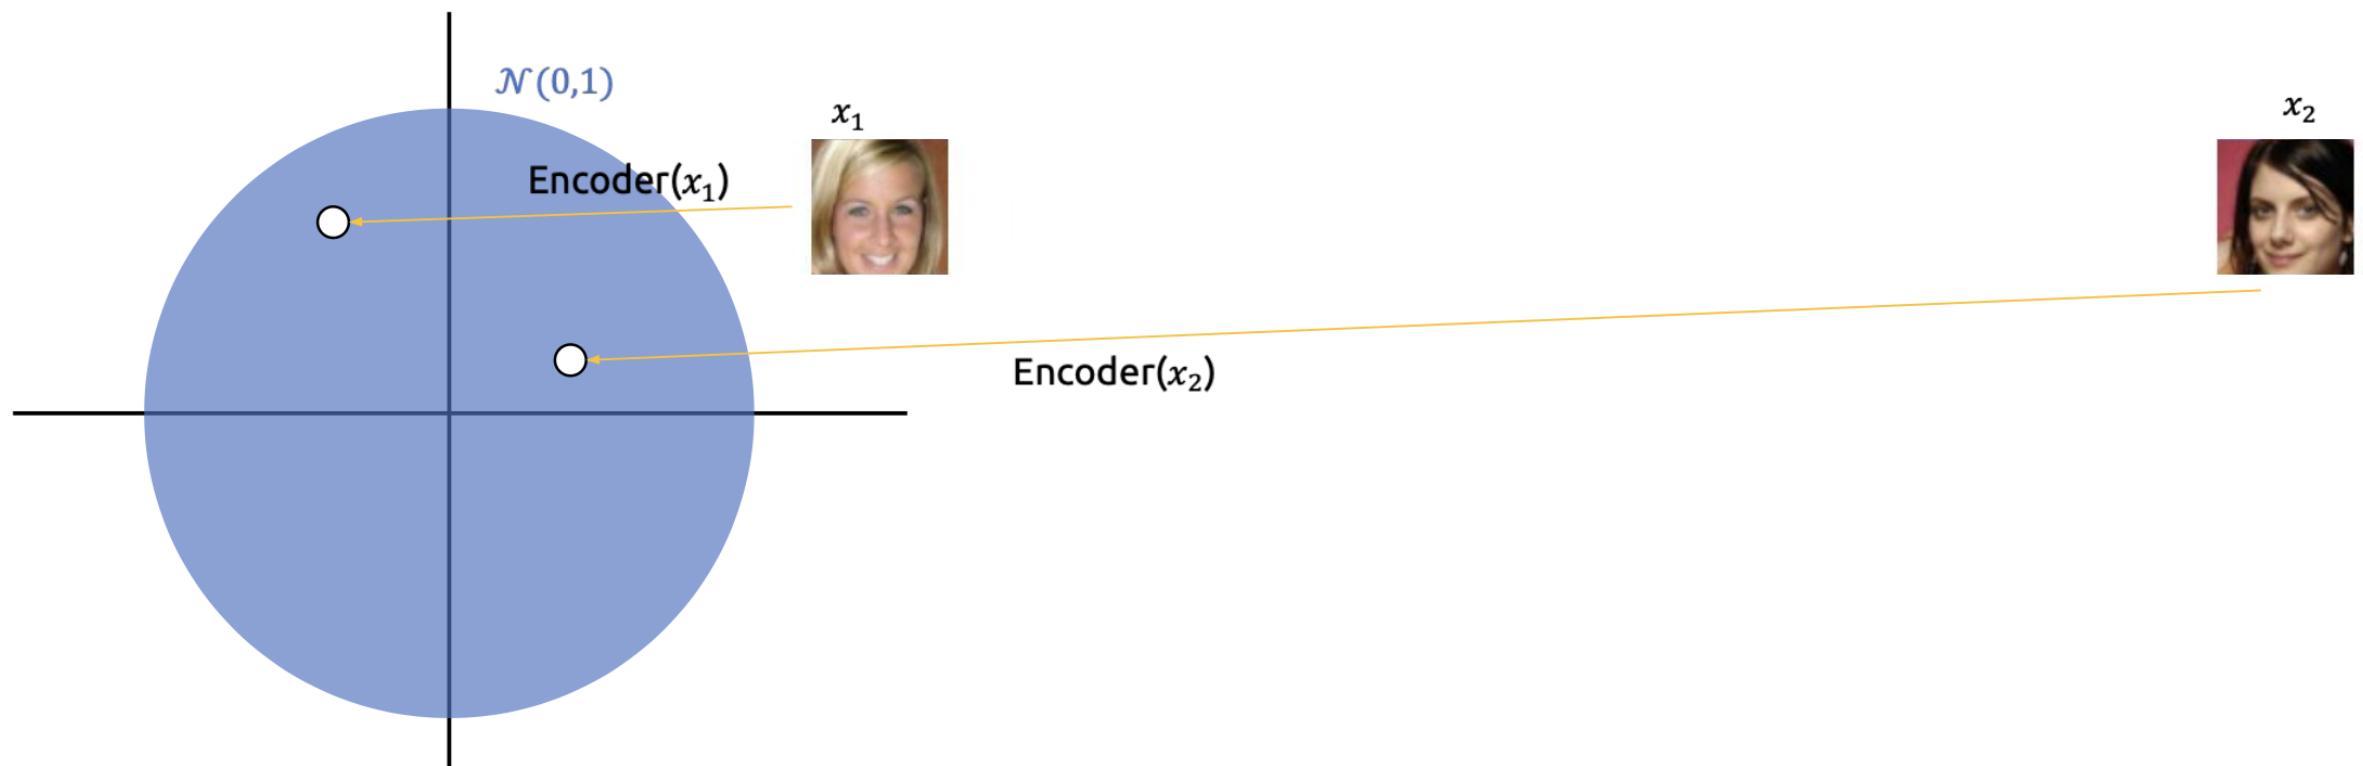
# Latent Space Interpolation

- Trace a linear path between two points in latent space, put all points along the path into the decoder



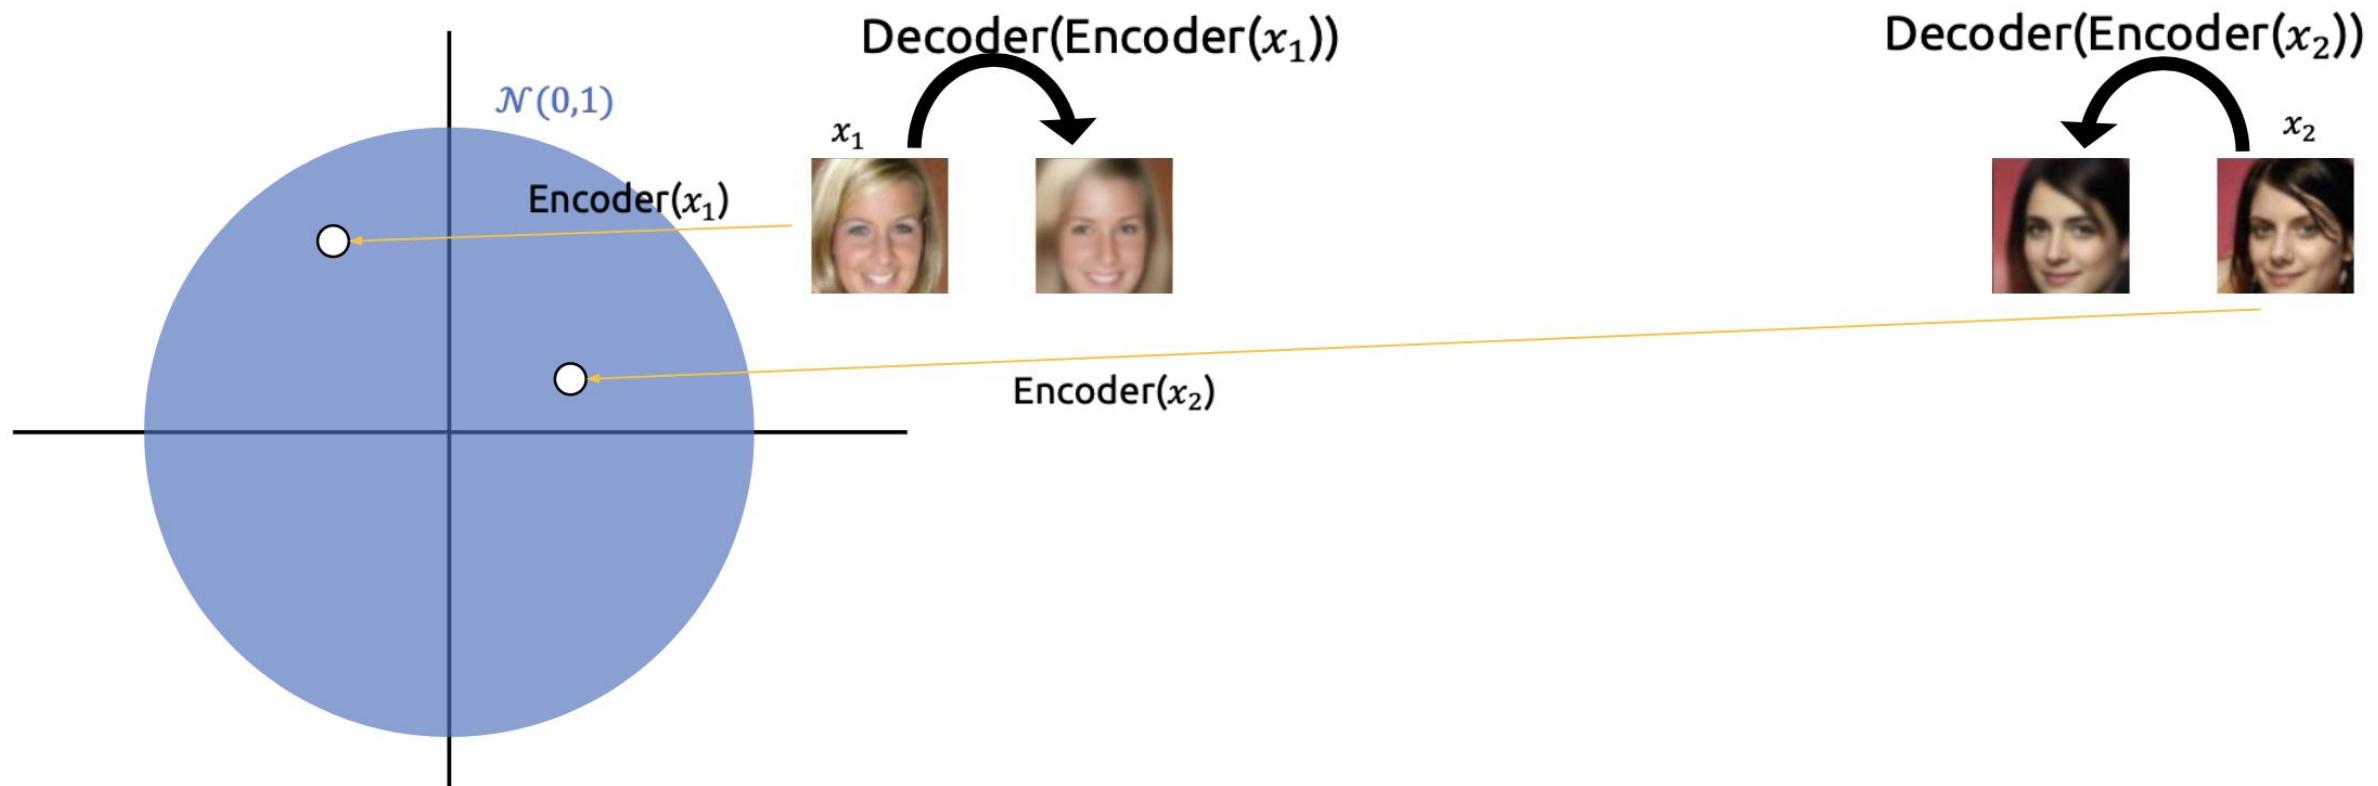
# Latent Space Interpolation

- Trace a linear path between two points in latent space, put all points along the path into the decoder



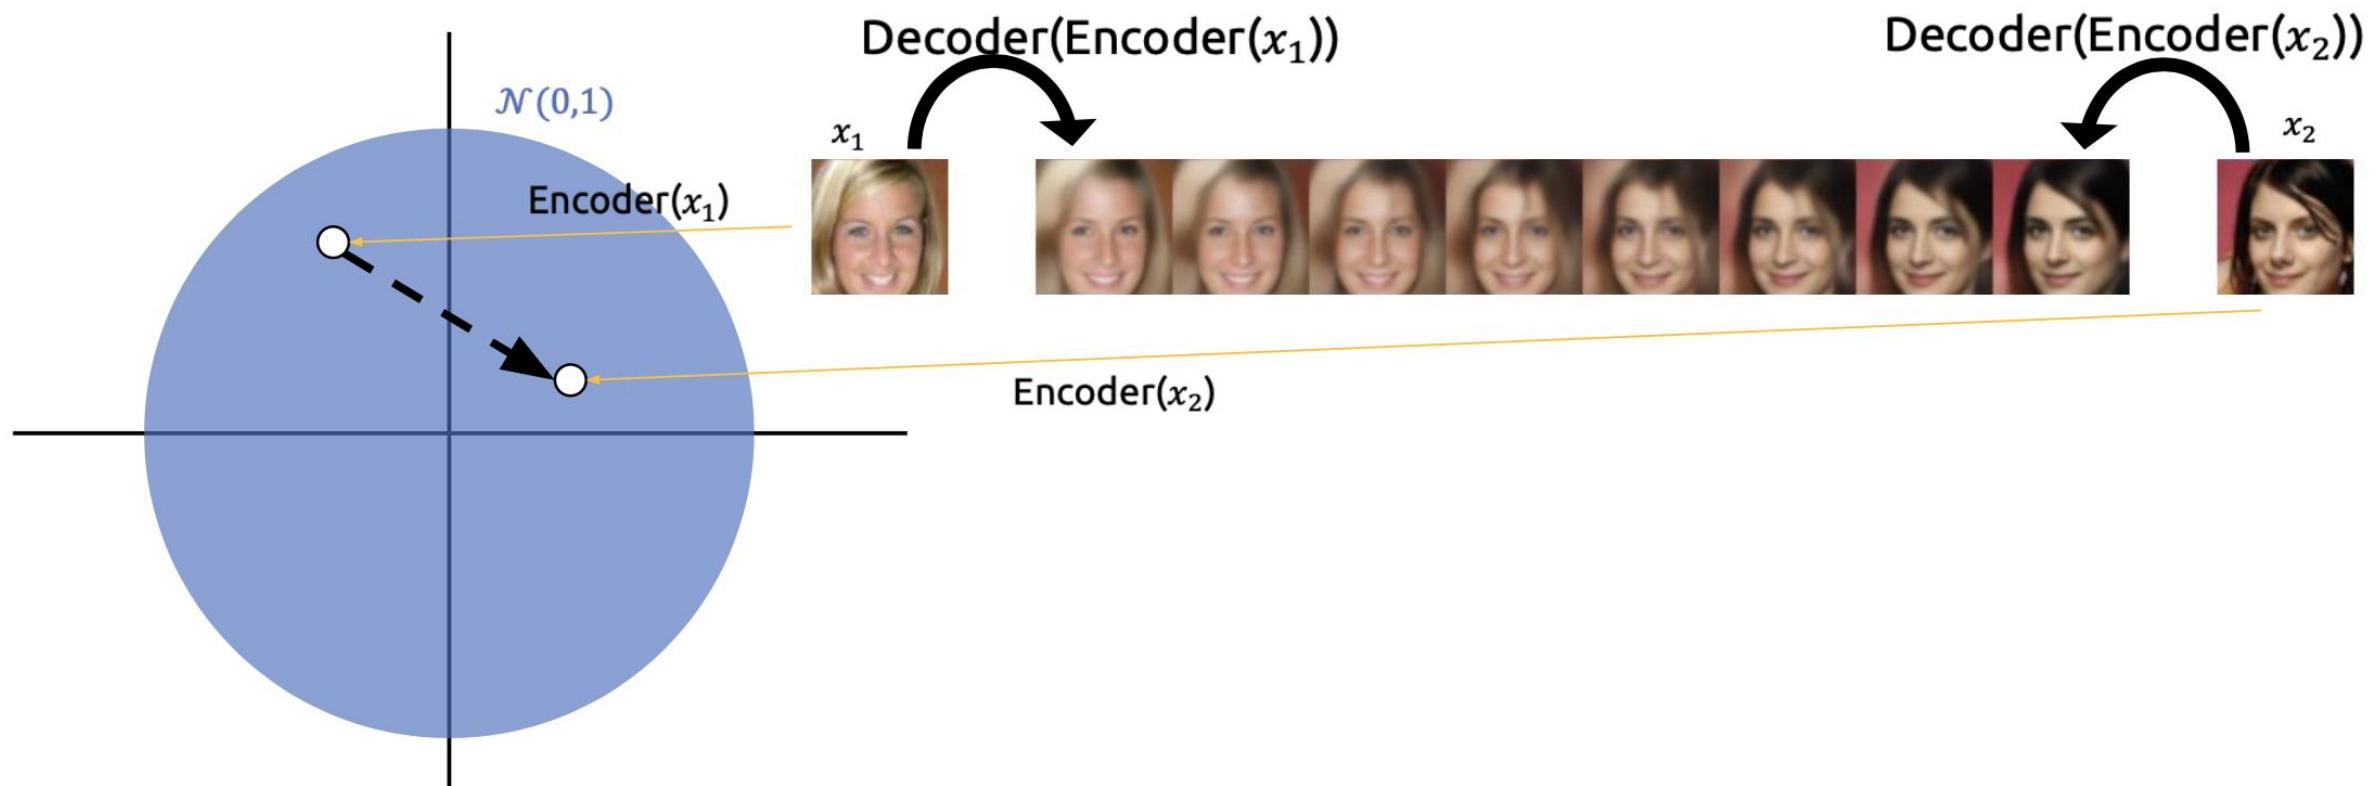
# Latent Space Interpolation

- Trace a linear path between two points in latent space, put all points along the path into the decoder



# Latent Space Interpolation

- Trace a linear path between two points in latent space, put all points along the path into the decoder



# Discriminative vs Generative Models

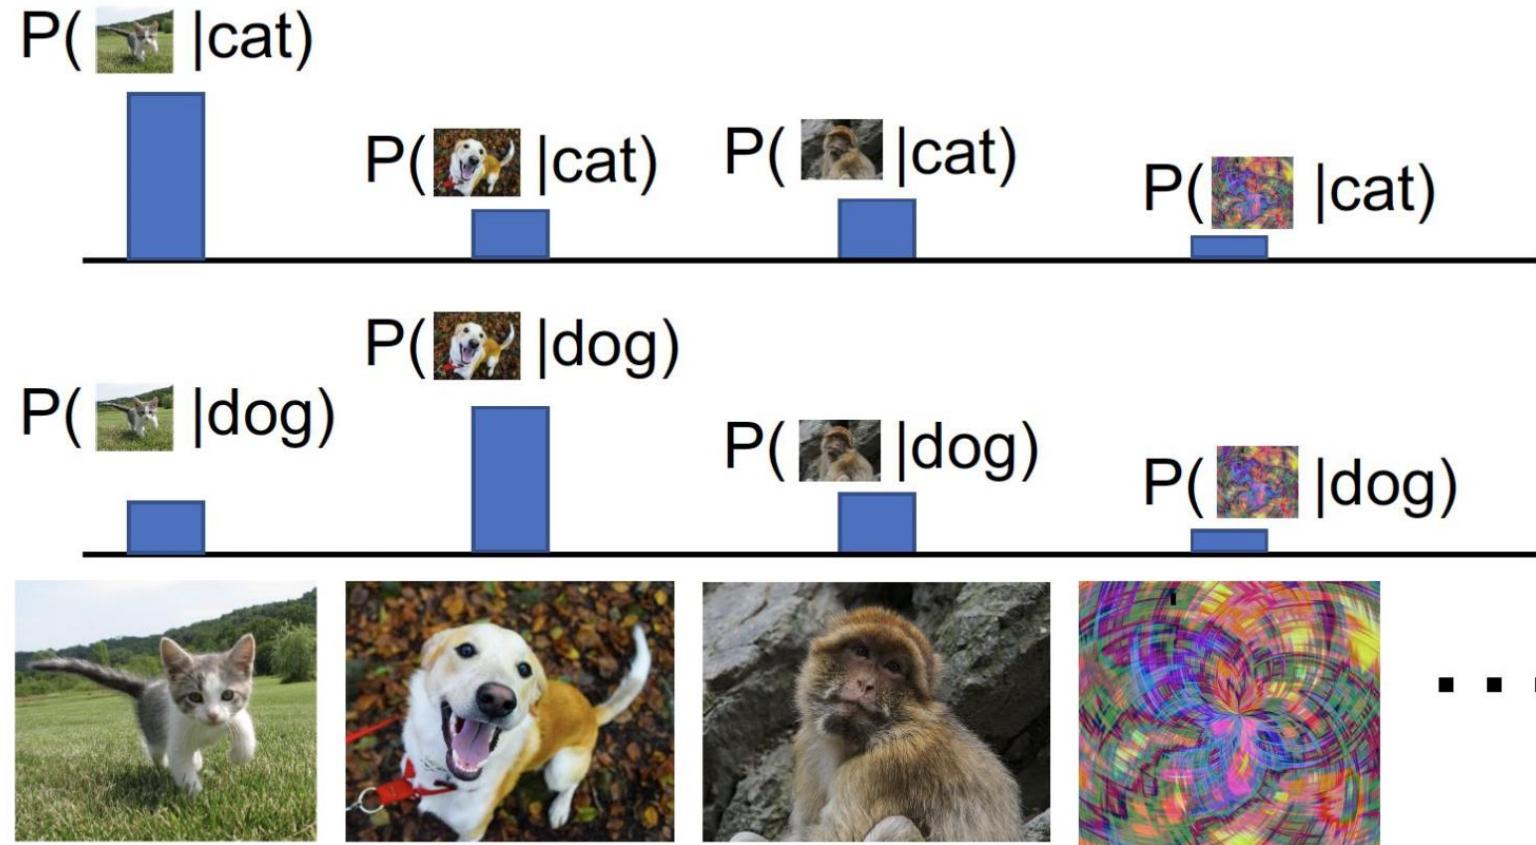
Discriminative Model:

Learn a probability distribution  $p(y|x)$

Generative Model:

Learn a probability distribution  $p(x)$

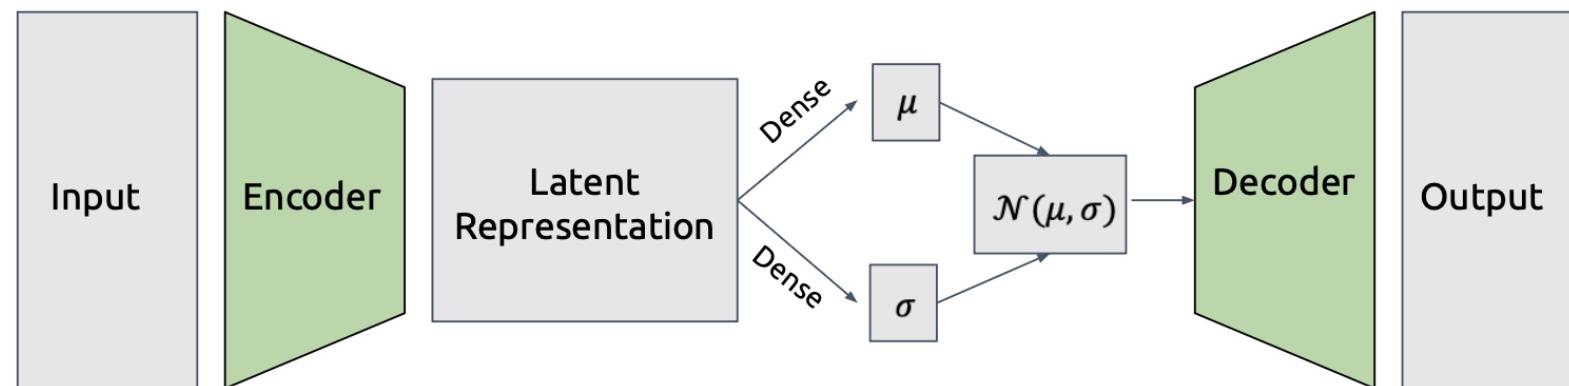
**Conditional Generative Model:** Learn  $p(x|y)$



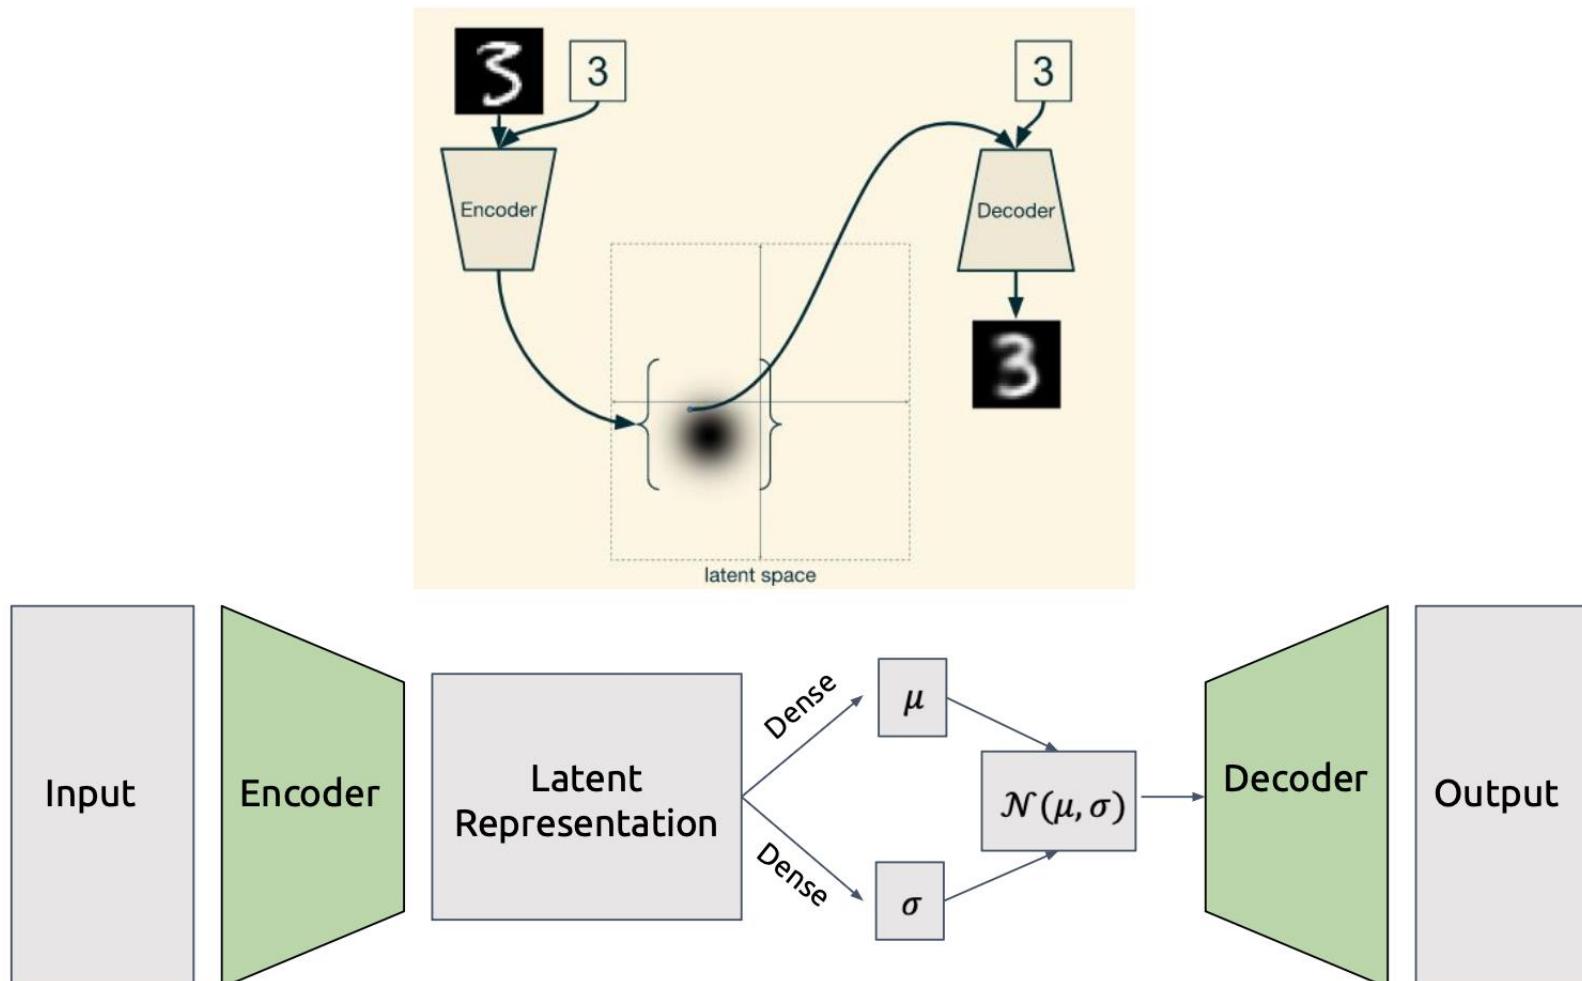
Conditional Generative Model: Each possible label induces a competition among all images

Any ideas?

# Conditional VAE



# Conditional VAE



# VAE output

Input



VAE reconstruction



What's the issue here?

Why?

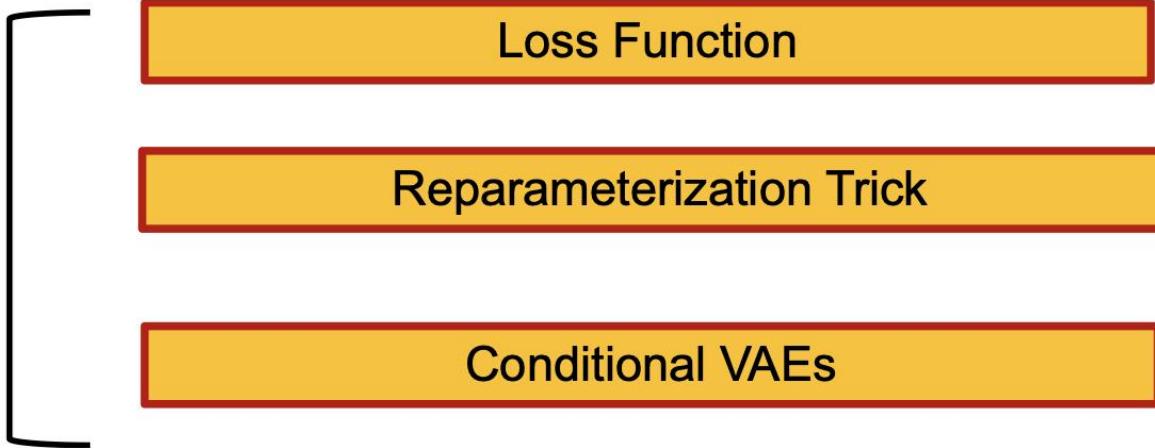
# Why are VAE samples blurry?

- Our reconstruction loss is the culprit
  - Mean Square Error (MSE) loss looks at each pixel in isolation
  - If no pixel is too far from its target value, the loss won't be too bad
  - Individual pixels look OK, but larger-scale features in the image aren't recognizable
- 
- ***Solutions?***
    - Let's choose a different reconstruction loss!



# Recap

Variational  
Autoencoders  
(VAEs)



<https://towardsdatascience.com/what-the-heck-are-vae-gans-17b86023588a>

